

Attention Mechanism

Part 1

Zhengyang Wang

Some slides are adapted from Stanford CS224N/Ling284.
(Abigail See and Richard Socher)


Outlines

- Review the attention mechanism in seq2seq models
- Definition of what the attention mechanism does
- How to use the attention mechanism in general


Outlines

- Review the attention mechanism in seq2seq models
- Definition of what the attention mechanism does
- How to use the attention mechanism in general


Sequence-to-sequence with attention



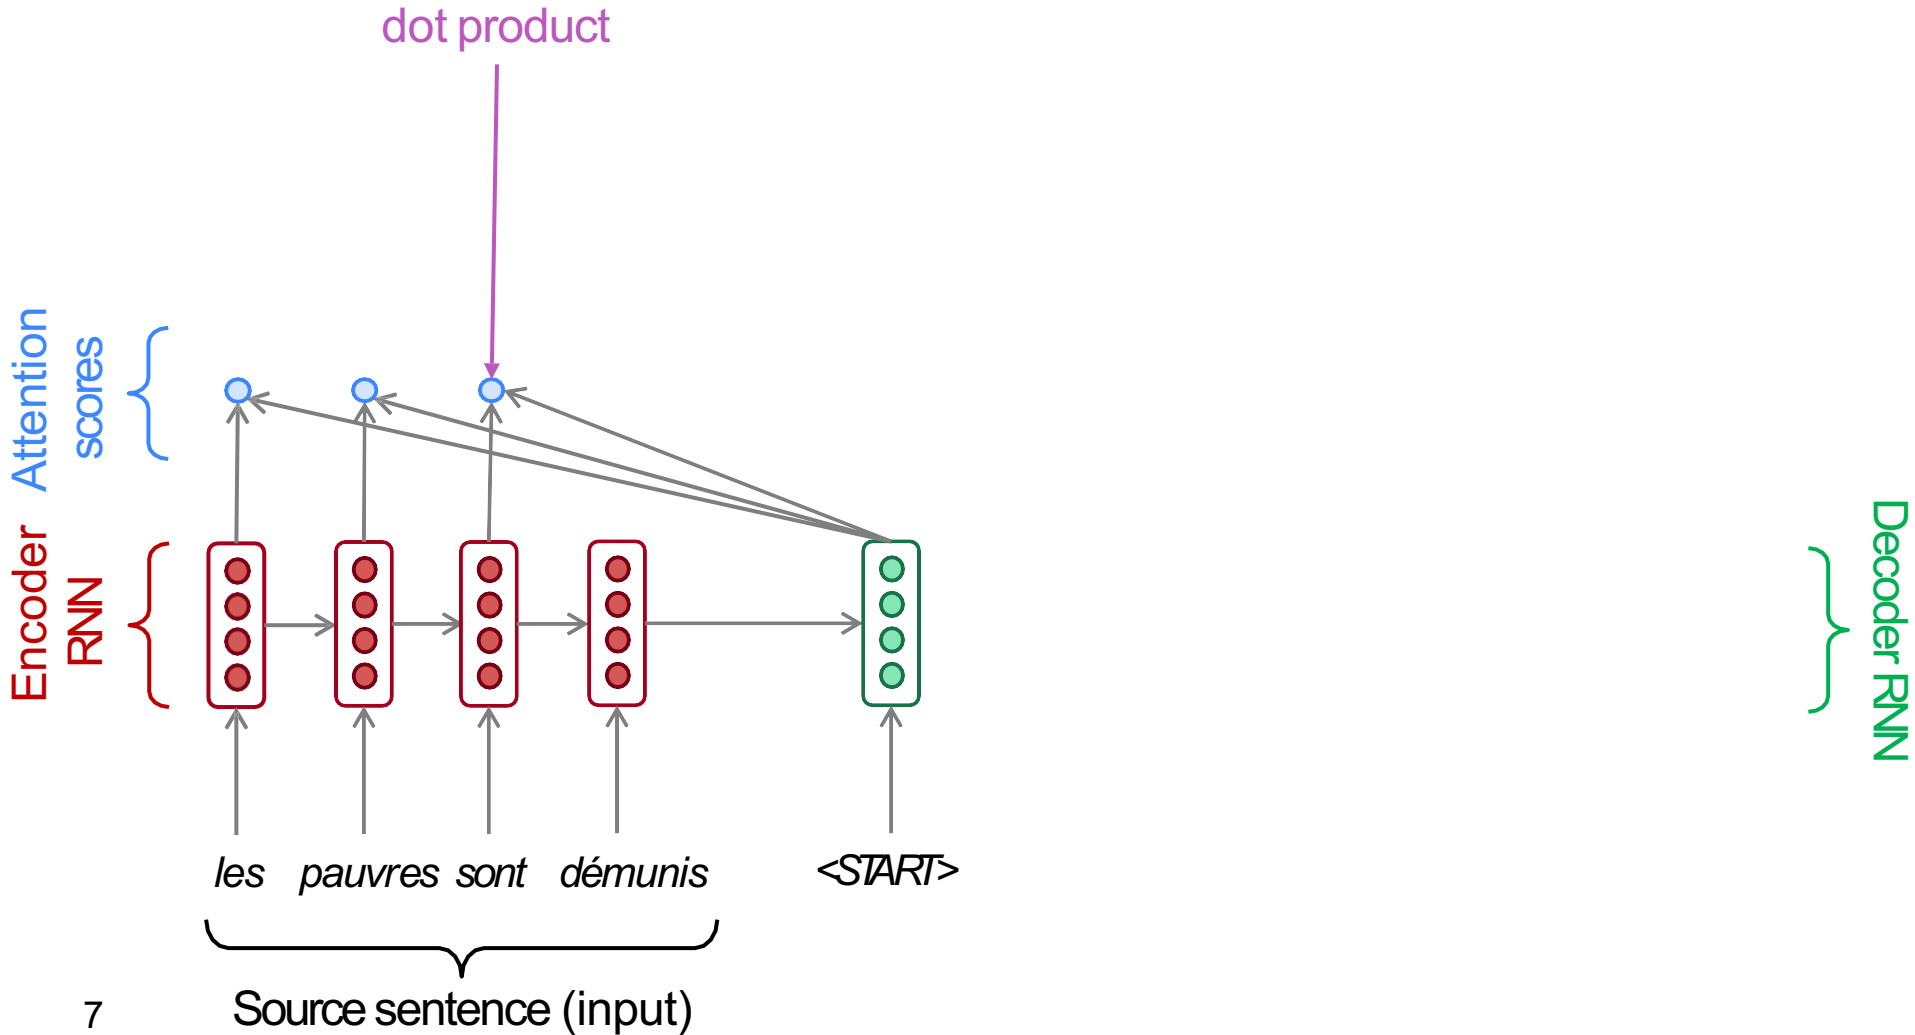
Sequence-to-sequence with attention




Sequence-to-sequence with attention



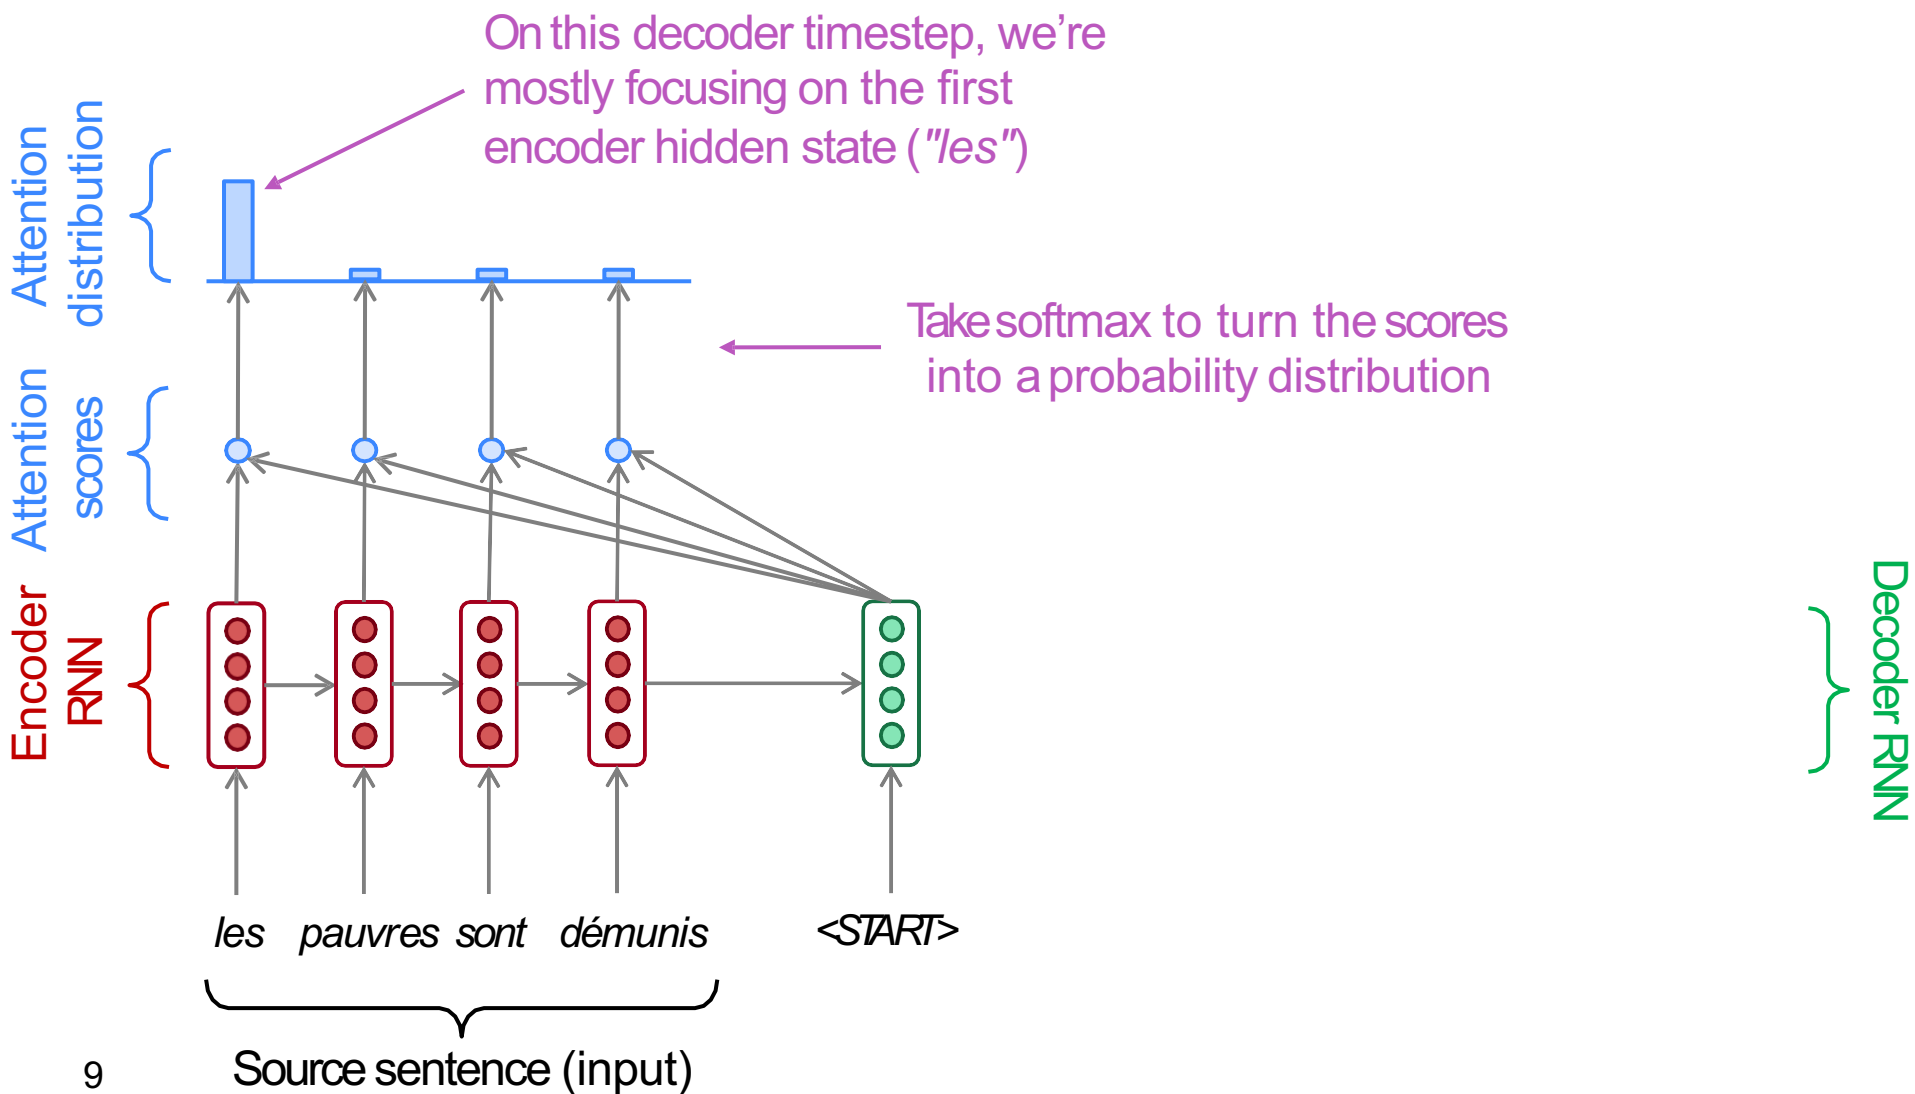
Sequence-to-sequence with attention



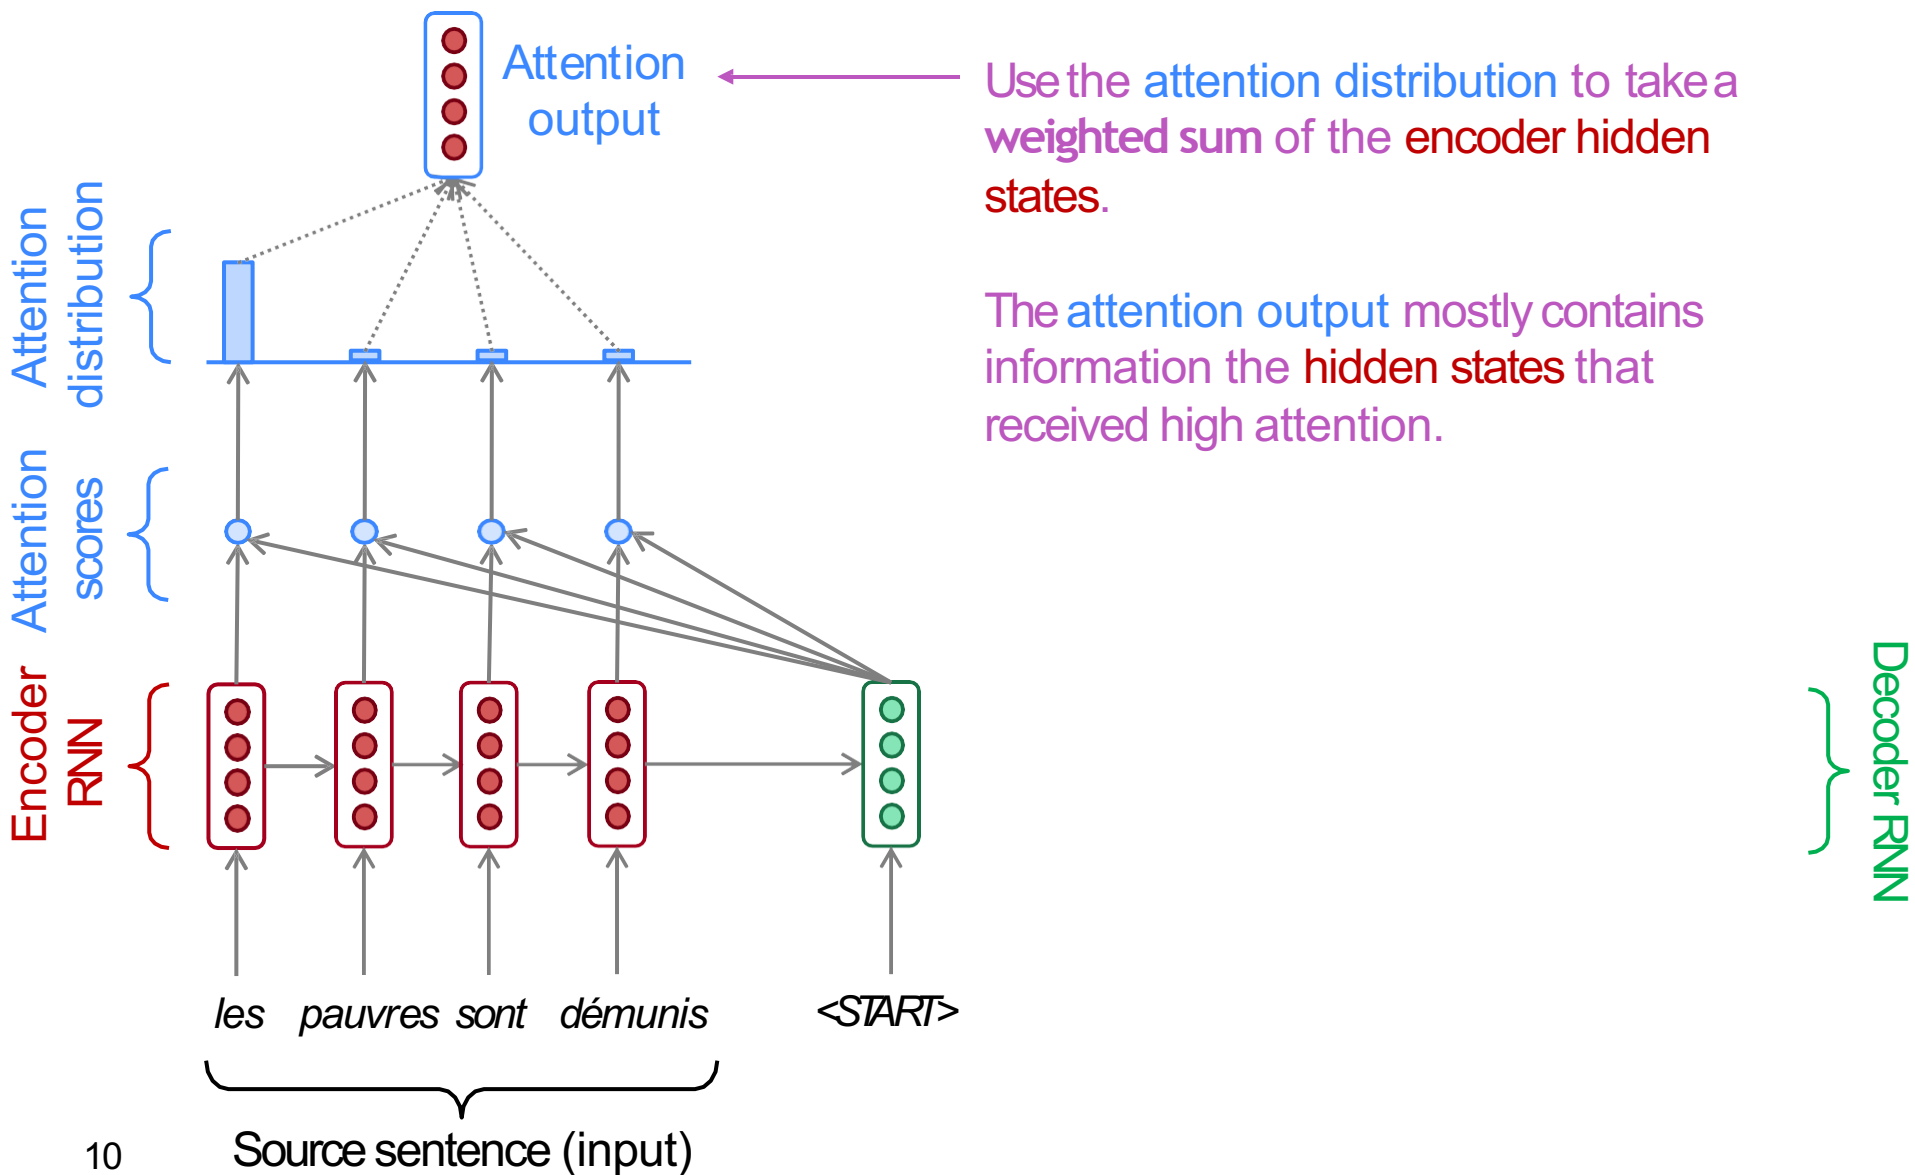
Sequence-to-sequence with attention



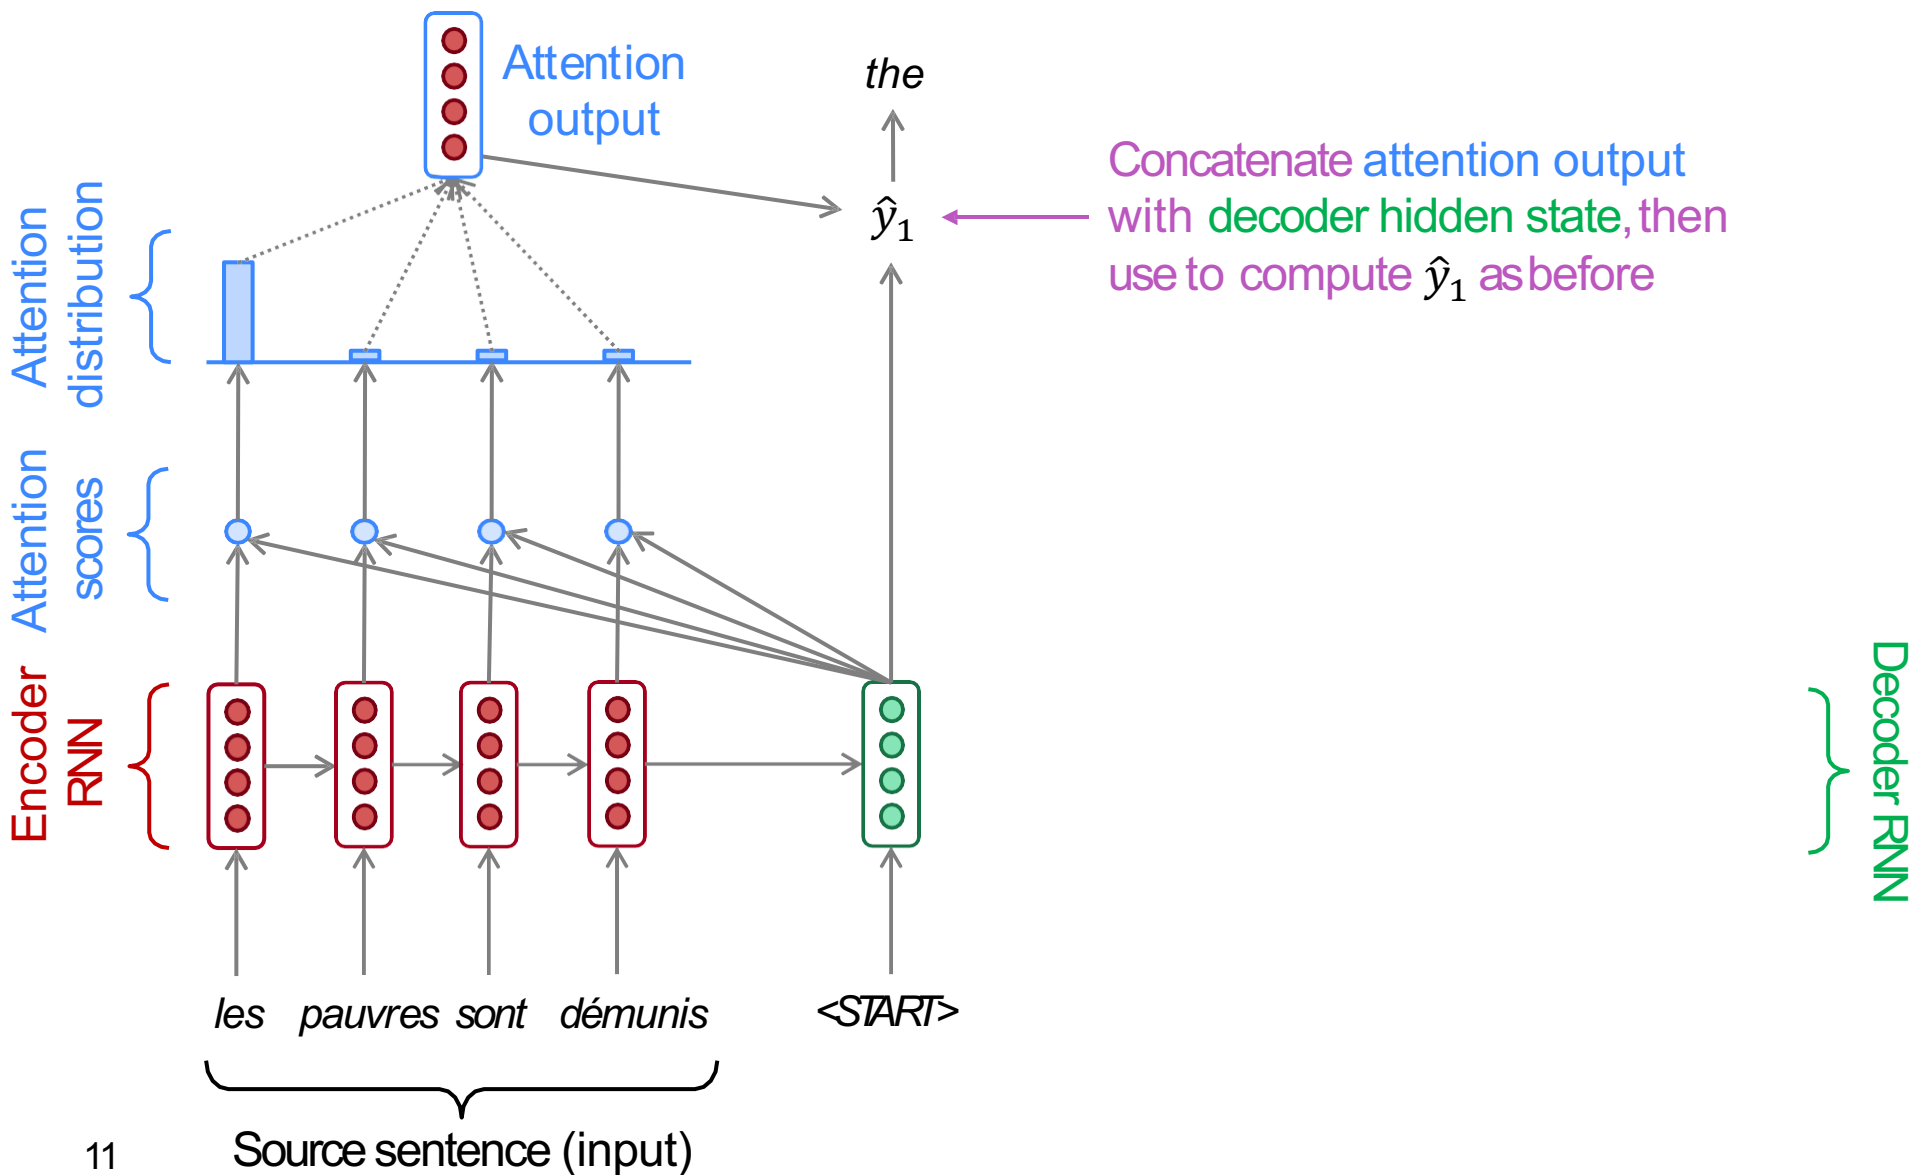
Sequence-to-sequence with attention



Sequence-to-sequence with attention



Sequence-to-sequence with attention



Attention: in equations

- We have encoder hidden states $\underline{h_1, \dots, h_N} \in \mathbb{R}^h$
- On timestep t , we have decoder hidden state $\underline{s_t} \in \mathbb{R}^h$
- We get the attention scores $\underline{e^t}$ for this step:

$$\underline{e^t} = [\underline{s_t^T h_1}, \dots, \underline{s_t^T h_N}] \in \mathbb{R}^N$$

- We take softmax to get the attention distribution $\underline{\alpha^t}$ for this step (this is a probability distribution and sums to 1)

$$\underline{\alpha^t} = \text{softmax}(\underline{e^t}) \in \mathbb{R}^N$$

- We use $\underline{\alpha^t}$ to take a weighted sum of the encoder hidden states to get the attention output $\underline{a_t}$

$$\underline{a_t} = \sum_{i=1}^N \alpha_i^t \underline{h_i} \in \mathbb{R}^h$$

- Finally we concatenate the attention output $\underline{a_t}$ with the decoder hidden state $\underline{s_t}$ and proceed as in the non-attention seq2seq model

$$[\underline{a_t}; \underline{s_t}] \in \mathbb{R}^{2h}$$

Outlines

- Review the attention mechanism in seq2seq models
- Definition of what the attention mechanism does
- How to use the attention mechanism in general

Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$ **Input 1**
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$ **Input 2**
- We get the attention scores e^t for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$ **Input 1**
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$ **Input 2**
- We get the attention scores e^t for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

Compute attention scores:
dot product

- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$ **Input 1**
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$ **Input 2**
- We get the attention scores e^t for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

Compute attention scores:
dot product

- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

Normalize attention scores:
softmax

- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$ **Input 1**
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$ **Input 2**
- We get the attention scores e^t for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

Compute attention scores:
dot product

- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

Normalize attention scores:
softmax

- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

Output

- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$ **Input 1**
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$ **Input 2**
- We get the attention scores e^t for this step:
$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$
 Compute attention scores: dot product
- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)
$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$
 Normalize attention scores: softmax
- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t
$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$
 Output: sum of Input 3 weighted by normalized attention scores
Input 3 ←
- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention: in equations

- We have encoder hidden states $h_1, \dots, h_N \in \mathbb{R}^h$ **Input-1 Key vectors**
- On timestep t , we have decoder hidden state $s_t \in \mathbb{R}^h$ **Input-2 Query vector(s)**
- We get the attention scores e^t for this step:

$$e^t = [s_t^T h_1, \dots, s_t^T h_N] \in \mathbb{R}^N$$

Compute attention scores:
dot product

- We take softmax to get the attention distribution α^t for this step (this is a probability distribution and sums to 1)

$$\alpha^t = \text{softmax}(e^t) \in \mathbb{R}^N$$

Normalize attention scores:
softmax

- We use α^t to take a weighted sum of the encoder hidden states to get the attention output a_t

$$a_t = \sum_{i=1}^N \alpha_i^t h_i \in \mathbb{R}^h$$

Value vectors **Input-3** ←

**Output: sum of Input 3
weighted by normalized
attention scores**

- Finally we concatenate the attention output a_t with the decoder hidden state s_t and proceed as in the non-attention seq2seq model

$$[a_t; s_t] \in \mathbb{R}^{2h}$$

Attention: definition

- Three inputs:
 - Query vector(s)
 - Key vectors
 - Value vectors
- Computation steps:
 - Compute attention scores
 - Normalize attention scores
- Output:
 - Sum of value vectors weighted by normalized attention scores

Attention: definition

- Three inputs:
 - Query vector(s)
 - Key vectors
 - Value vectors



- Three input matrices:
 - $Q \in \mathbb{R}^{d_q \times n_q}$
 - $K \in \mathbb{R}^{d_k \times n_k}$
 - $V \in \mathbb{R}^{d_v \times n_v}$


**We have $n_q = 1$ in the seq2seq example.
We stick to $n_q = 1$ for now.**

Attention: definition


- Three inputs:
 - Query vector(s)
 - Key vectors
 - Value vectors
- Three input matrices:
 - $Q \in \mathbb{R}^{d_q \times n_q}$
 - $K \in \mathbb{R}^{d_k \times n_k}$
 - $V \in \mathbb{R}^{d_v \times n_v}$
- Computation steps:
 - Compute attention scores: compute the dot product between Q ($n_q = 1$) and each column in K .

$$A = Q^T K \in \mathbb{R}^{n_q \times n_k}$$


Attention: definition

- Three inputs:
 - Query vector(s)
 - Key vectors
 - Value vectors
 - Three input matrices:
 - $Q \in \mathbb{R}^{d_q \times n_q}$
 - $K \in \mathbb{R}^{d_k \times n_k}$
 - $V \in \mathbb{R}^{d_v \times n_v}$
- 
- Computation steps:
 - Compute attention scores: compute the dot product between Q ($n_q = 1$) and each column in K .
- $$A = Q^T K \in \mathbb{R}^{n_q \times n_k}$$
- Does it put any constraint on the first dimension of Q and K ?

Attention: definition

- Three inputs:
 - Query vector(s)
 - Key vectors
 - Value vectors
 - Three input matrices:
 - $Q \in \mathbb{R}^{d_q \times n_q}$
 - $K \in \mathbb{R}^{d_k \times n_k}$
 - $V \in \mathbb{R}^{d_v \times n_v}$
- 
- Computation steps:
 - Compute attention scores: compute the dot product between Q ($n_q = 1$) and each column in K .
- $$A = Q^T K \in \mathbb{R}^{n_q \times n_k}$$
- Does it put any constraint on the first dimension of Q and K ?
$$d_q = d_k$$

Attention: definition

- Three inputs:
 - Query vector(s)
 - Key vectors
 - Value vectors
- 
- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:


$$d_q = d_k$$

- Computation steps:
 - Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
 - **Normalize attention scores: Softmax**
 $A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: Softmax

$$A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$$

- If $n_q \neq 1$, should we perform Softmax over column or row?

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: Softmax

$$A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$$

- If $n_q \neq 1$, should we perform Softmax over column or row?

Row

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: $A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$

- Output:

- Sum of value vectors weighted by normalized attention scores:

$$\text{Output} = V \cdot A^T \in \mathbb{R}^{d_v \times n_q}$$

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: $A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$

- Output:

- Sum of value vectors weighted by normalized attention scores:

$$\text{Output} = V \cdot A^T \in \mathbb{R}^{d_v \times n_q}$$

- Does it put any constraint on the second dimension of V and K ?

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: $A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$

- Output:

- Sum of value vectors weighted by normalized attention scores:

$$\text{Output} = V \cdot A^T \in \mathbb{R}^{d_v \times n_q}$$

- Does it put any constraint on the second dimension of V and K ?

$$n_v = n_k$$

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

$$n_v = n_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: $A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$

- Output:


- Sum of value vectors weighted by normalized attention scores:

$$\text{Output} = V \cdot A^T \in \mathbb{R}^{d_v \times n_q}$$

Attention: definition

- Three inputs:

- Query vector(s)
- Key vectors
- Value vectors



- Three input matrices:

- $Q \in \mathbb{R}^{d_q \times n_q}$
- $K \in \mathbb{R}^{d_k \times n_k}$
- $V \in \mathbb{R}^{d_v \times n_v}$

Constraints:

$$d_q = d_k$$

$$n_v = n_k$$

- Computation steps:

- Compute attention scores: $A = Q^T K \in \mathbb{R}^{n_q \times n_k}$
- Normalize attention scores: $A = \text{Softmax}(A) \in \mathbb{R}^{n_q \times n_k}$

- Output:

- Sum of value vectors weighted by normalized attention scores:

$$\text{Output} = V \cdot A^T \in \mathbb{R}^{d_v \times n_q}$$

- The shape of the output is determined by the number of query vectors and the dimension of value vectors.

Outlines

- Review the attention mechanism in seq2seq models
- Definition of what the attention mechanism does
- How to use the attention mechanism in general

Use Attention

- Use attention:
 - Determine Q, K, V
 - In the seq2seq example, we have $K = V$ (encoder hidden states) and a single query vector $Q = q$ (the decoder hidden state at time t).

Use Attention

- Use attention:
 - Determine Q, K, V
 - In the seq2seq example, we have $K = V$ (encoder hidden states) and a single query vector $Q = q$ (the decoder hidden state at time t).
 - Define the way to compute and normalize attention scores
 - Dot product + Softmax is just one way.
 - Check paper “Non-local Neural Networks” for other ways.
 - For example: dot product + $1/n_k$; use a neural network.

Use Attention

- Can we use attention just as convolution or fully-connected layer?
 - One input X
 - Some training parameters
 - One output Y
- How can we have Q, K, V if we only have one input X ?

Use Attention

- Can we use attention just as convolution or fully-connected layer?
 - One input X
 - Some training parameters
 - One output Y
- How can we have Q, K, V if we only have one input X ?
- Self-Attention
 - $Q = K = V = X$
 - Any problem?

Use Attention

- Can we use attention just as convolution or fully-connected layer?
 - One input X
 - Some training parameters
 - One output Y
- How can we have Q, K, V if we only have one input X ?
- Self-Attention
 - $Q = K = V = X$
 - Any problem?
 - No training parameter involved.
 - It does not make much sense.

Self-Attention

- Instead of having $Q = K = V = X$, use X to generate Q, K, V .
 - Suppose $X \in \mathbb{R}^{d_x \times n_x}$
 - Three independent linear transformations:
 - $Q = W_q X \in \mathbb{R}^{d_q \times n_x}$
 - $K = W_k X \in \mathbb{R}^{d_k \times n_x}$
 - $V = W_v X \in \mathbb{R}^{d_v \times n_x}$
 - $W_q \in \mathbb{R}^{d_q \times d_x}, W_k \in \mathbb{R}^{d_k \times d_x}, W_v \in \mathbb{R}^{d_v \times d_x}$ are trainable parameters.
 - W_q, W_k must satisfy the constraints $d_q = d_k$.

Self-Attention

- Instead of having $Q = K = V = X$, use X to generate Q, K, V .
 - Suppose $X \in \mathbb{R}^{d_x \times n_x}$
 - Three independent linear transformations:
 - $Q = W_q X \in \mathbb{R}^{d_q \times n_x}$
 - $K = W_k X \in \mathbb{R}^{d_k \times n_x}$
 - $V = W_v X \in \mathbb{R}^{d_v \times n_x}$
 - $W_q \in \mathbb{R}^{d_q \times d_x}, W_k \in \mathbb{R}^{d_k \times d_x}, W_v \in \mathbb{R}^{d_v \times d_x}$ are trainable parameters.
 - W_q, W_k must satisfy the constraints $d_q = d_k$.
- What is the shape of the output Y ?

Self-Attention

- Instead of having $Q = K = V = X$, use X to generate Q, K, V .
 - Suppose $X \in \mathbb{R}^{d_x \times n_x}$
 - Three independent linear transformations:
 - $Q = W_q X \in \mathbb{R}^{d_q \times n_x}$
 - $K = W_k X \in \mathbb{R}^{d_k \times n_x}$
 - $V = W_v X \in \mathbb{R}^{d_v \times n_x}$
 - $W_q \in \mathbb{R}^{d_q \times d_x}, W_k \in \mathbb{R}^{d_k \times d_x}, W_v \in \mathbb{R}^{d_v \times d_x}$ are trainable parameters.
 - W_q, W_k must satisfy the constraints $d_q = d_k$.
- What is the shape of the output Y ?
 $d_v \times n_x$

Self-Attention

- Instead of having $Q = K = V = X$, use X to generate Q, K, V .
 - Suppose $X \in \mathbb{R}^{d_x \times n_x}$
 - Three independent linear transformations:
 - $Q = W_q X \in \mathbb{R}^{d_q \times n_x}$
 - $K = W_k X \in \mathbb{R}^{d_k \times n_x}$
 - $V = W_v X \in \mathbb{R}^{d_v \times n_x}$
 - $W_q \in \mathbb{R}^{d_q \times d_x}, W_k \in \mathbb{R}^{d_k \times d_x}, W_v \in \mathbb{R}^{d_v \times d_x}$ are trainable parameters.
 - W_q, W_k must satisfy the constraints $d_q = d_k$.
- What is the shape of the output Y ?
 $d_v \times n_x$
- If we omit W_v and simply have $V = X$, each column in Y is a weighted sum of all column vectors in X .

Self-Attention

- If we omit W_v and simply have $V = X$, each column in Y is a weighted sum of all column vectors in X .
- Comparison with convolution and fully-connected layer

Self-Attention

- If we omit W_v and simply have $V = X$, each column in Y is a weighted sum of all column vectors in X .
- Comparison with convolution and fully-connected layer
 - Convolution
 - Each column in Y is computed from column vectors within a local range in X .

Self-Attention

- If we omit W_v and simply have $V = X$, each column in Y is a weighted sum of all column vectors in X .
- Comparison with convolution and fully-connected layer
 - Convolution
 - Each column in Y is computed from column vectors within a local range in X .
 - Fully-connected layer
 - The weights in the weighted sum is not input-dependent.

Multi-Head Self-Attention

- Instead of having $Q = K = V = X$, use X to generate Q, K, V .
 - Suppose $X \in \mathbb{R}^{d_x \times n_x}$
 - Three independent linear transformations:
 - $Q = W_q X \in \mathbb{R}^{d_q \times n_x}$
 - $K = W_k X \in \mathbb{R}^{d_k \times n_x}$
 - $V = W_v X \in \mathbb{R}^{d_v \times n_x}$
 - $W_q \in \mathbb{R}^{d_q \times d_x}, W_k \in \mathbb{R}^{d_k \times d_x}, W_v \in \mathbb{R}^{d_v \times d_x}$ are trainable parameters.
 - W_q, W_k must satisfy the constraints $d_q = d_k$.
- Do the above process for multiple times independently.
 - Each time results in an output Y_i .
 - Concatenate all the Y_i as the final output.