

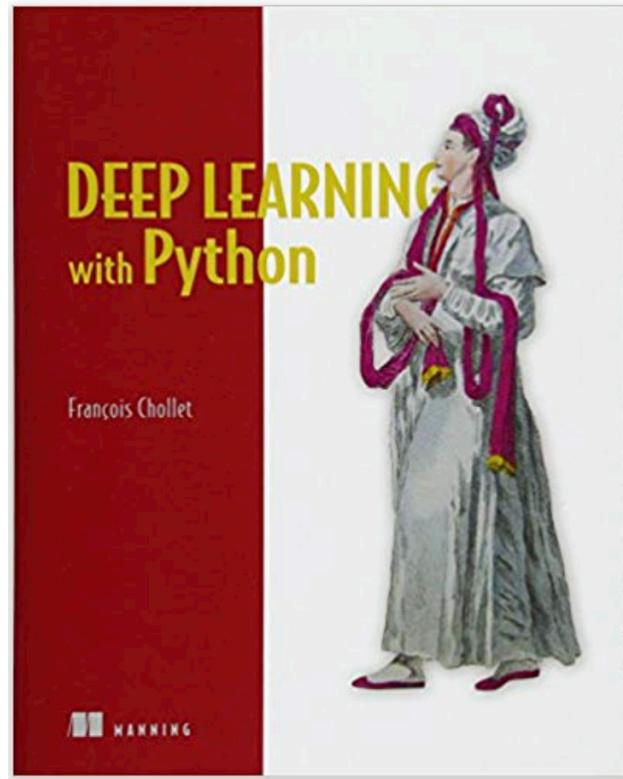
CSCE 636 Neural Networks (Deep Learning)

Lecture 2: Mathematical Building Blocks of Neural Networks

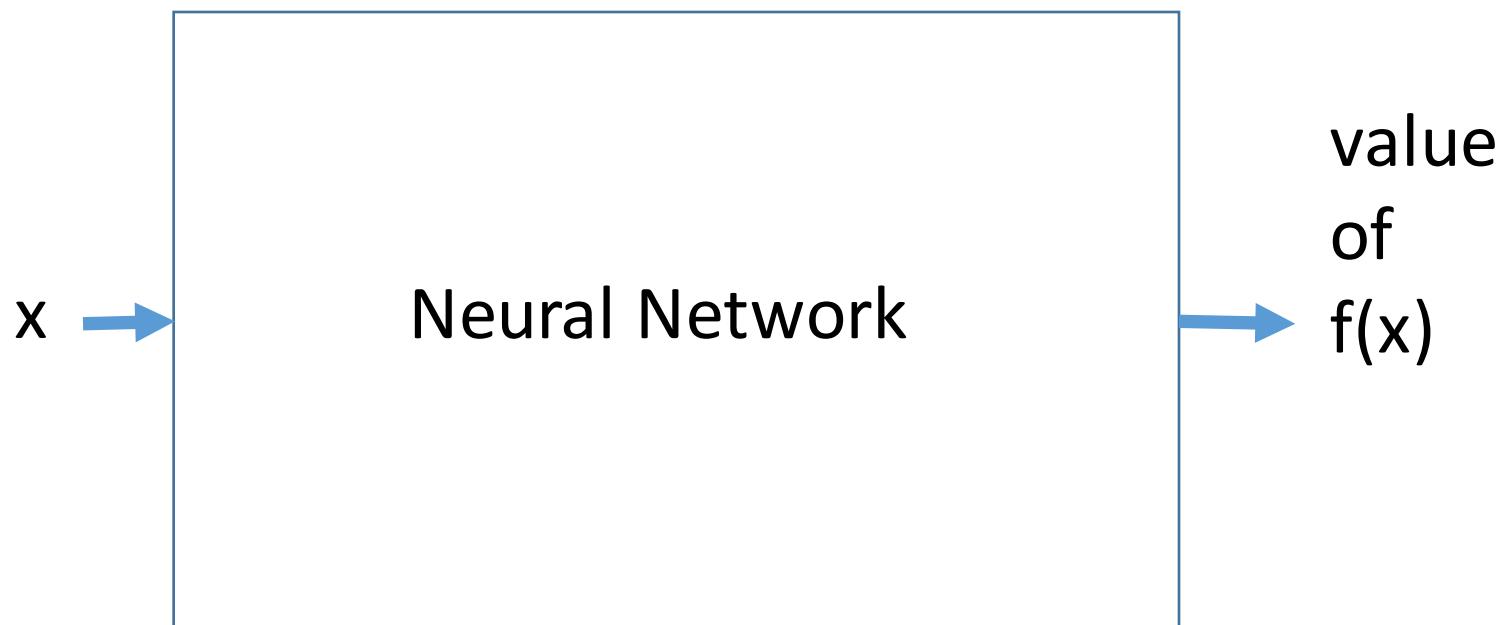
Anxiao (Andrew) Jiang

Chapter 2

Before we begin: the mathematical building blocks of neural networks



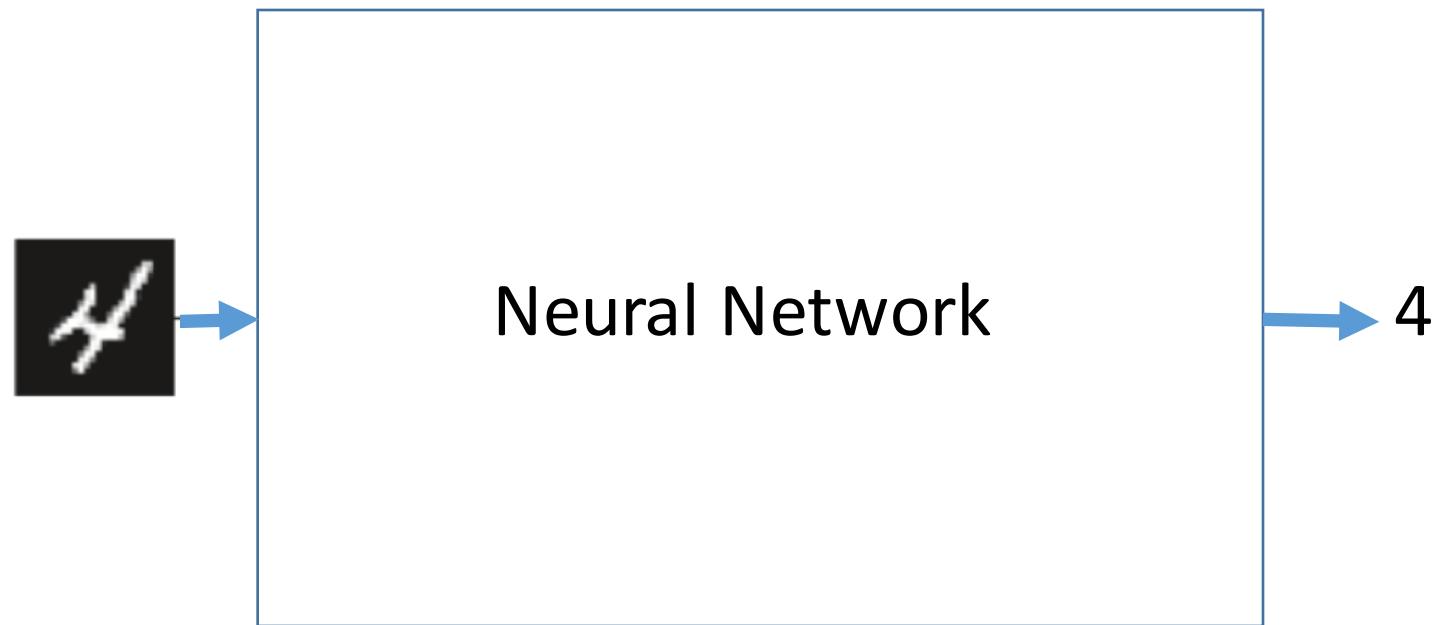
What a neural network does: learn a function



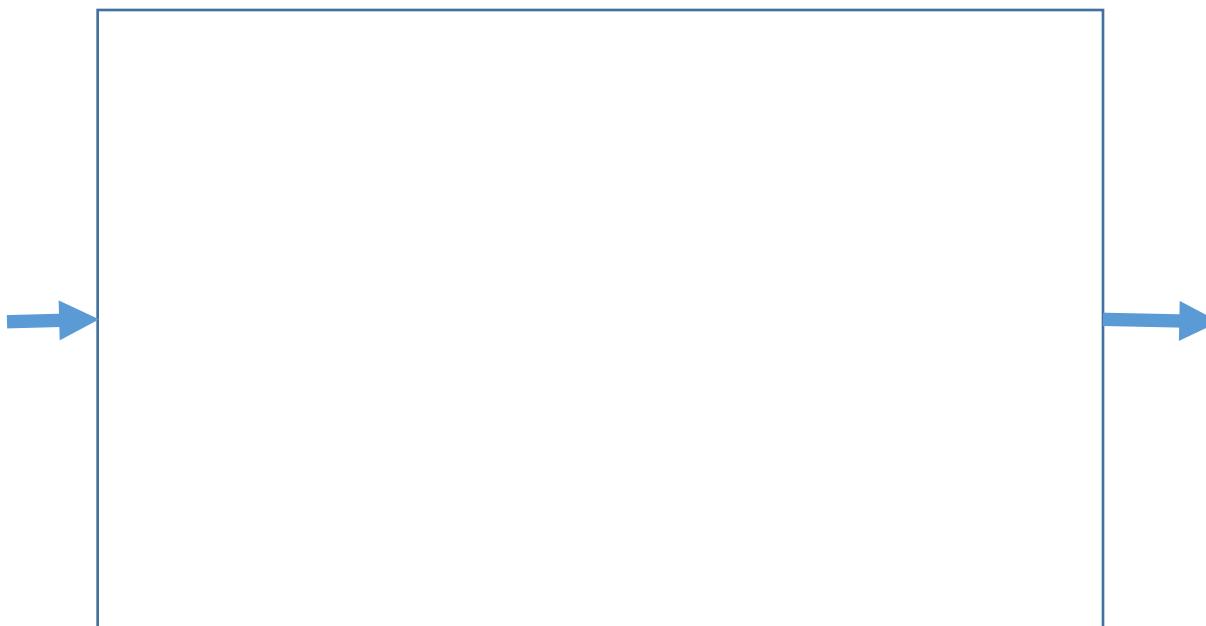
The neural network learns the function $f(x)$,
either exactly or approximately.

Application: Handwritten Digit Recognition

Task: Classify grayscale images of handwritten digits (28x28 pixels) into their 10 categories (0 through 9).



How to start?



Step 1: Load the dataset

MNIST Dataset: 60,000 training images and 10,000 test images, along with their labels.



Step 1: Load the dataset

Listing 2.1 Loading the MNIST dataset in Keras

```
from keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

train_images: 60,000 x 28 x 28 array, where each element (pixel) is an integer in [0,255]

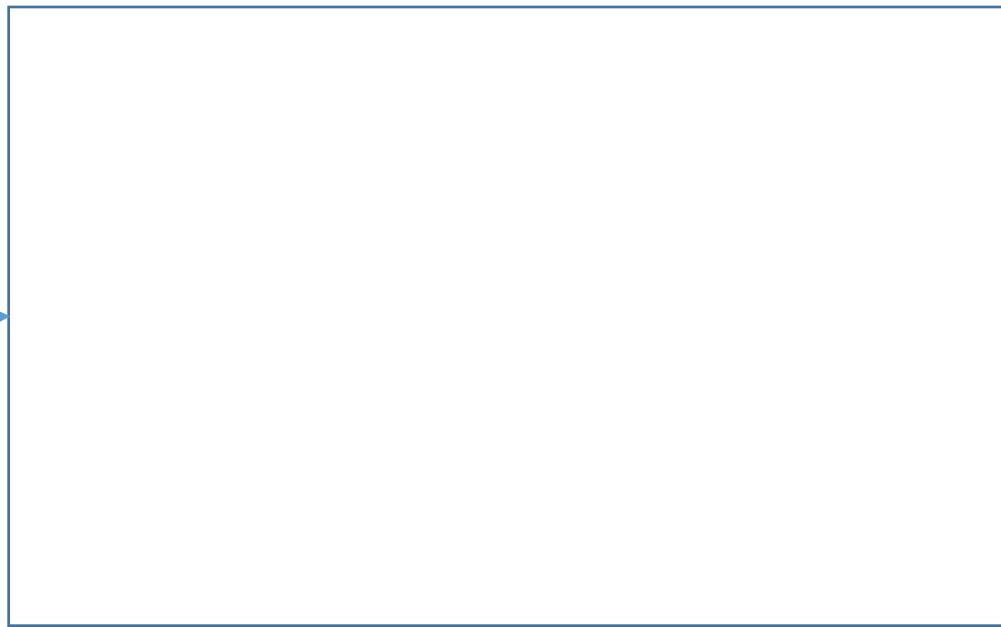
train_labels: vector of length 60,000, where each element (label) is an integer in [0,9]

test_images: 10,000 x 28 x 28 array, where each element (pixel) is an integer in [0,255]

test_labels: vector of length 10,000, where each element (label) is an integer in [0,9]

Rule: training data and test data are disjoint. Only use training data to train neural network!

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9



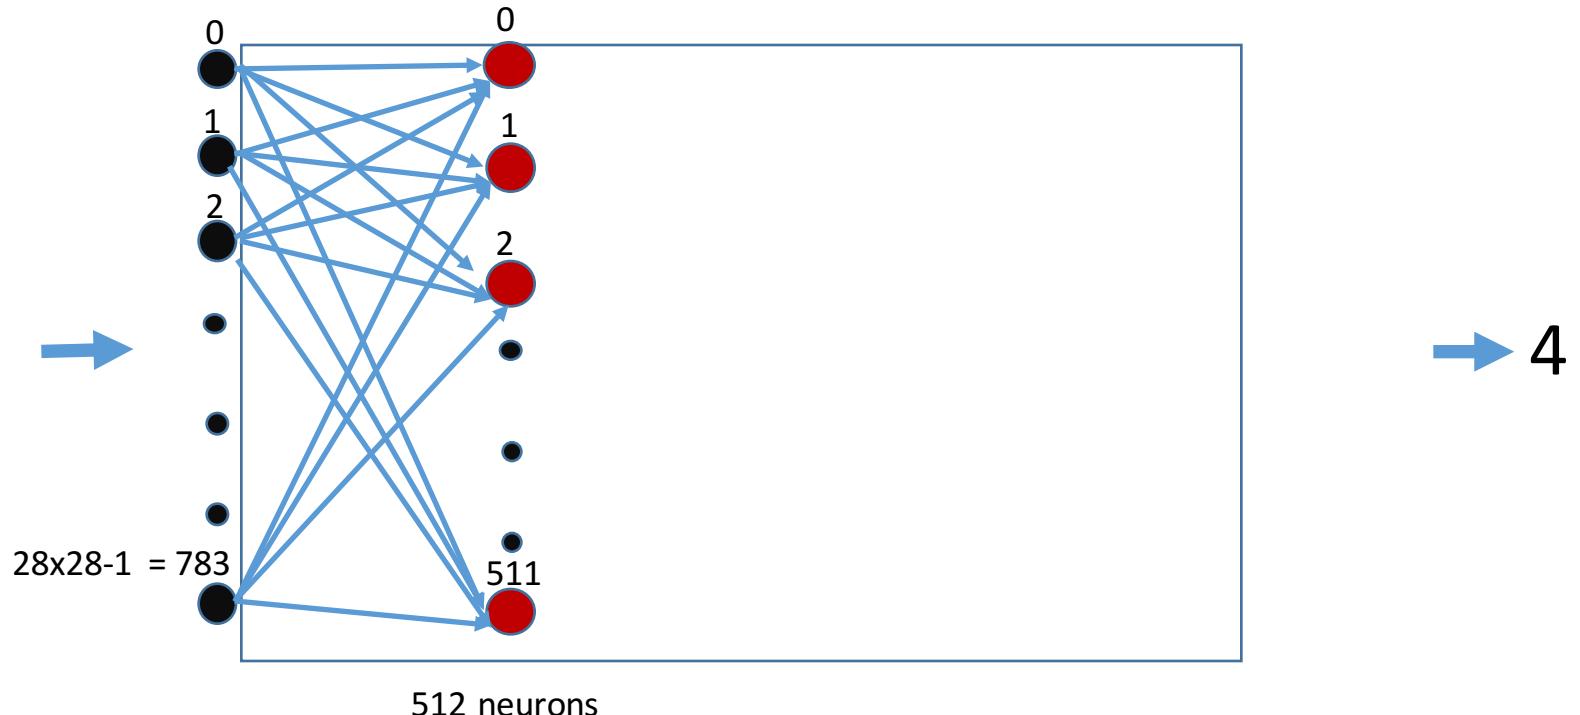
4

Step 2: Build neural network architecture

```
from keras import models  
from keras import layers  
  
network = models.Sequential()  
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
```

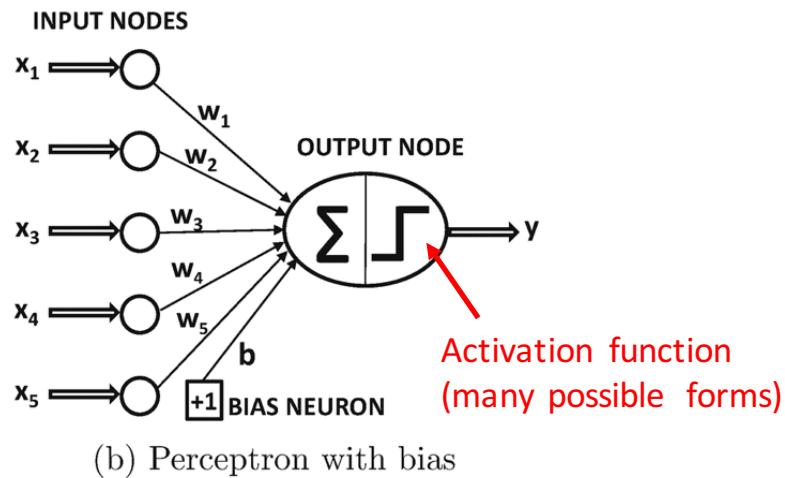


28 x 28
2-d array

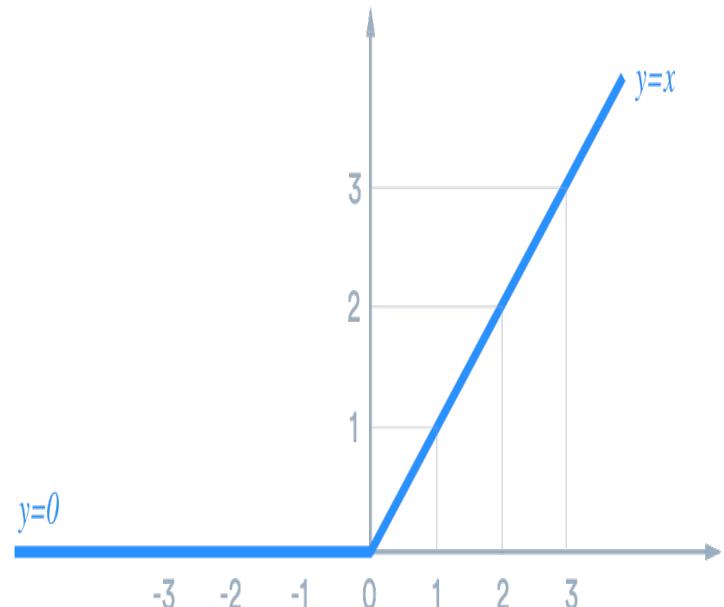


What is a neuron

ReLU (most popular Activation function)



$$\hat{y} = \text{sign}\{\overline{W} \cdot \overline{X} + b\} = \text{sign}\left\{\sum_{j=1}^d w_j x_j + b\right\}$$

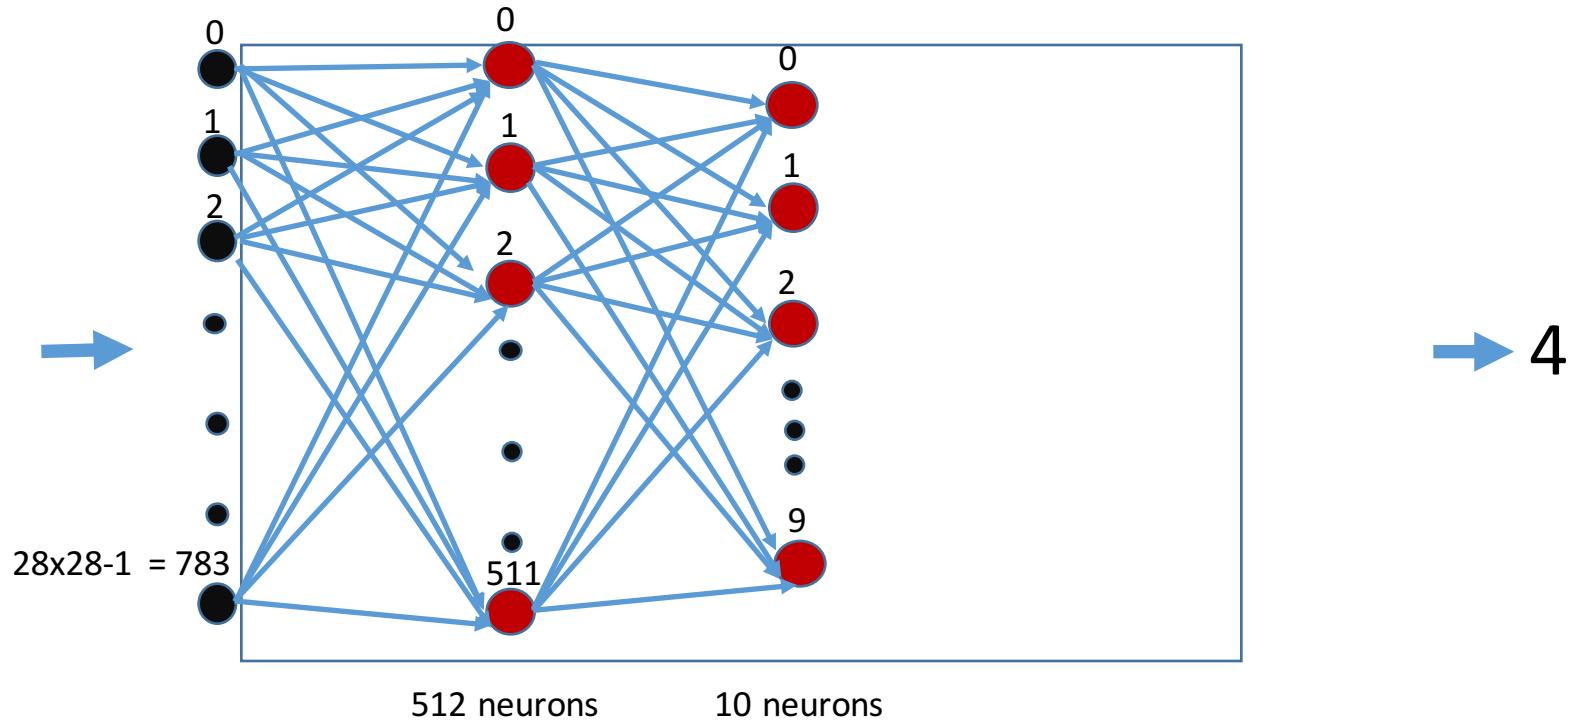


Step 2: Build neural network architecture

```
from keras import models  
from keras import layers  
  
network = models.Sequential()  
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))  
network.add(layers.Dense(10, activation='softmax'))
```



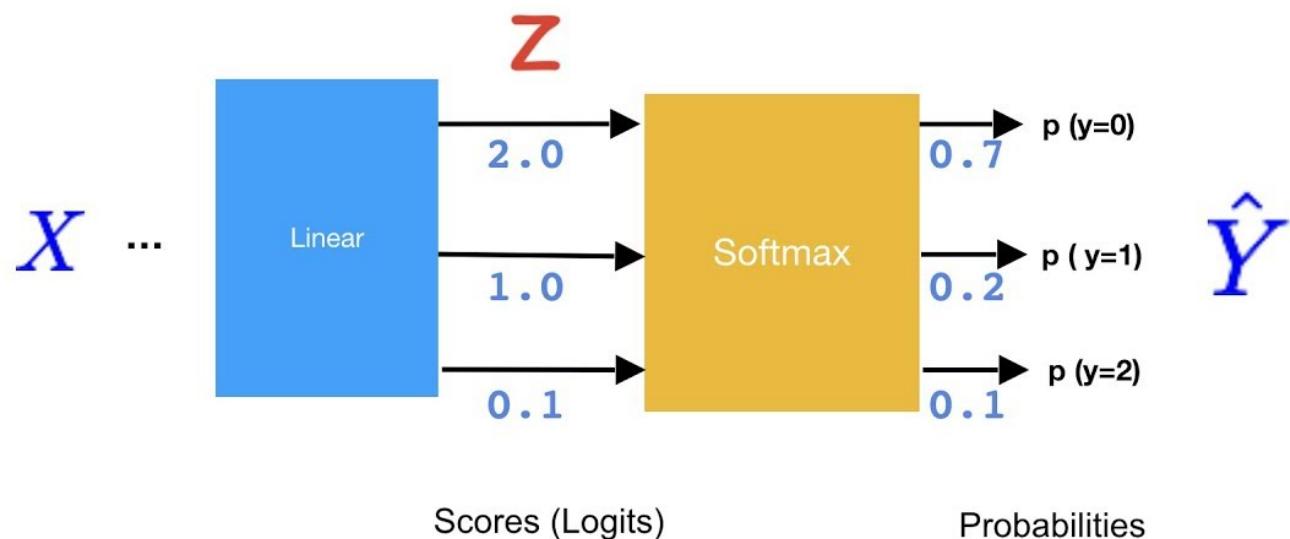
28 x 28
2-d array



Softmax (popular activation function for the last layer of a classification network)

Meet Softmax

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

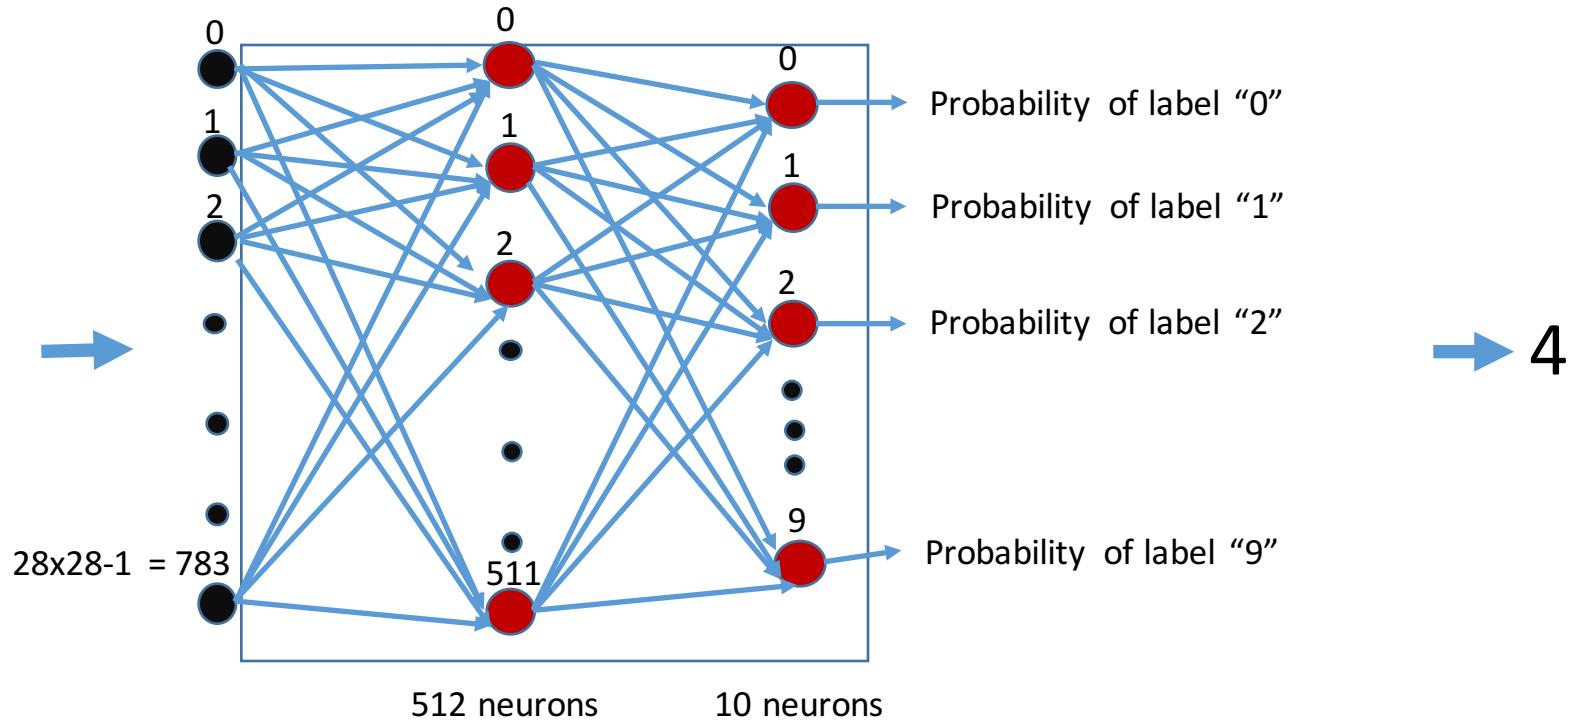


Step 2: Build neural network architecture

```
from keras import models  
from keras import layers  
  
network = models.Sequential()  
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))  
network.add(layers.Dense(10, activation='softmax'))
```



28 x 28
2-d array



Step 3: choose loss function, optimizer, and target metrics

Listing 2.3 The compilation step

```
network.compile(optimizer='rmsprop',  
                loss='categorical_crossentropy',  
                metrics=[ 'accuracy' ] )
```

Categorical cross-entropy (a popular loss function for multi-class classification)

$$CCE = -\frac{1}{N} \sum_{i=0}^N \sum_{j=0}^J y_j \cdot \log(\hat{y}_j) + (1 - y_j) \cdot \log(1 - \hat{y}_j)$$

Number of classes

Number of samples

True probability (1 or 0)
this input belongs to
class j

probability predicted by neural
network that this input belongs to
class j

RMSProp (a popular optimizer, details to be introduced later)

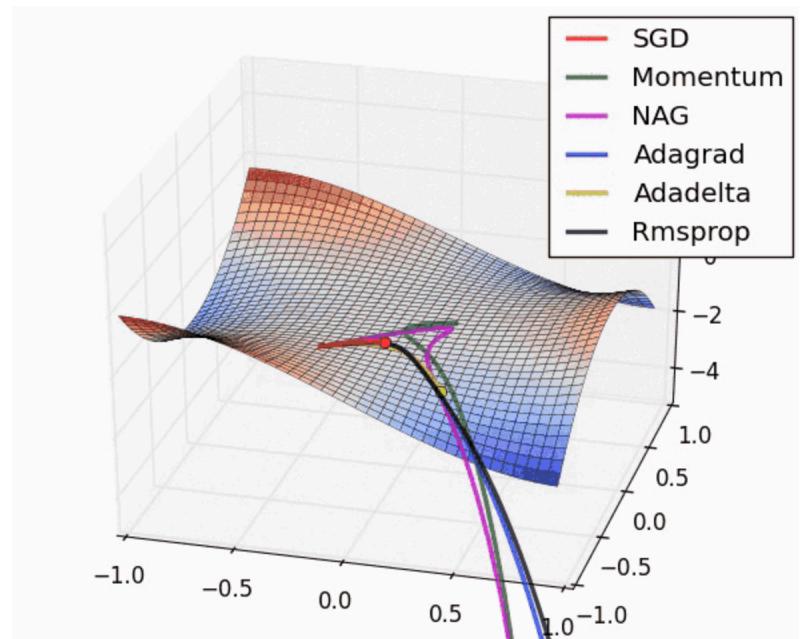
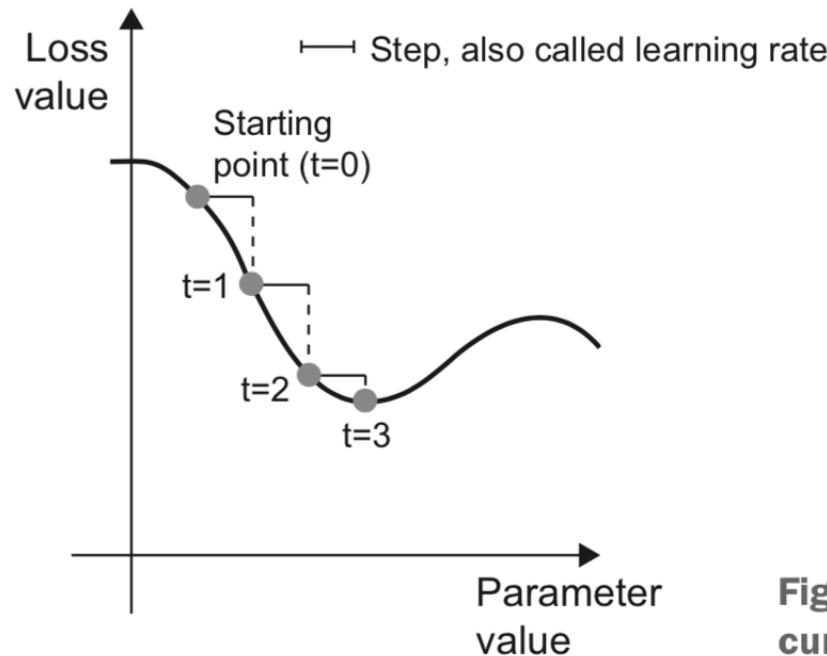
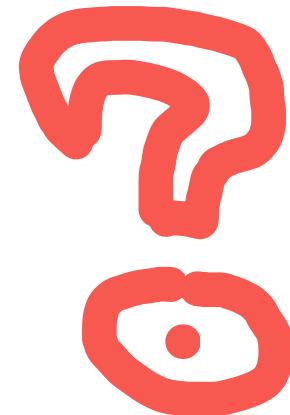


Figure 2.11 SGD down a 1D loss curve (one learnable parameter)

Accuracy: fraction of times that the neural network makes correct predictions

- If we care about accuracy, why do we optimize categorical cross-entropy during training?
- Answer: loss function needs to be differentiable. (And the loss function is closely related to the target metric. Minimizing the loss function is (approximately or precisely) optimizing the target metric.)





The “**Teacher**”:

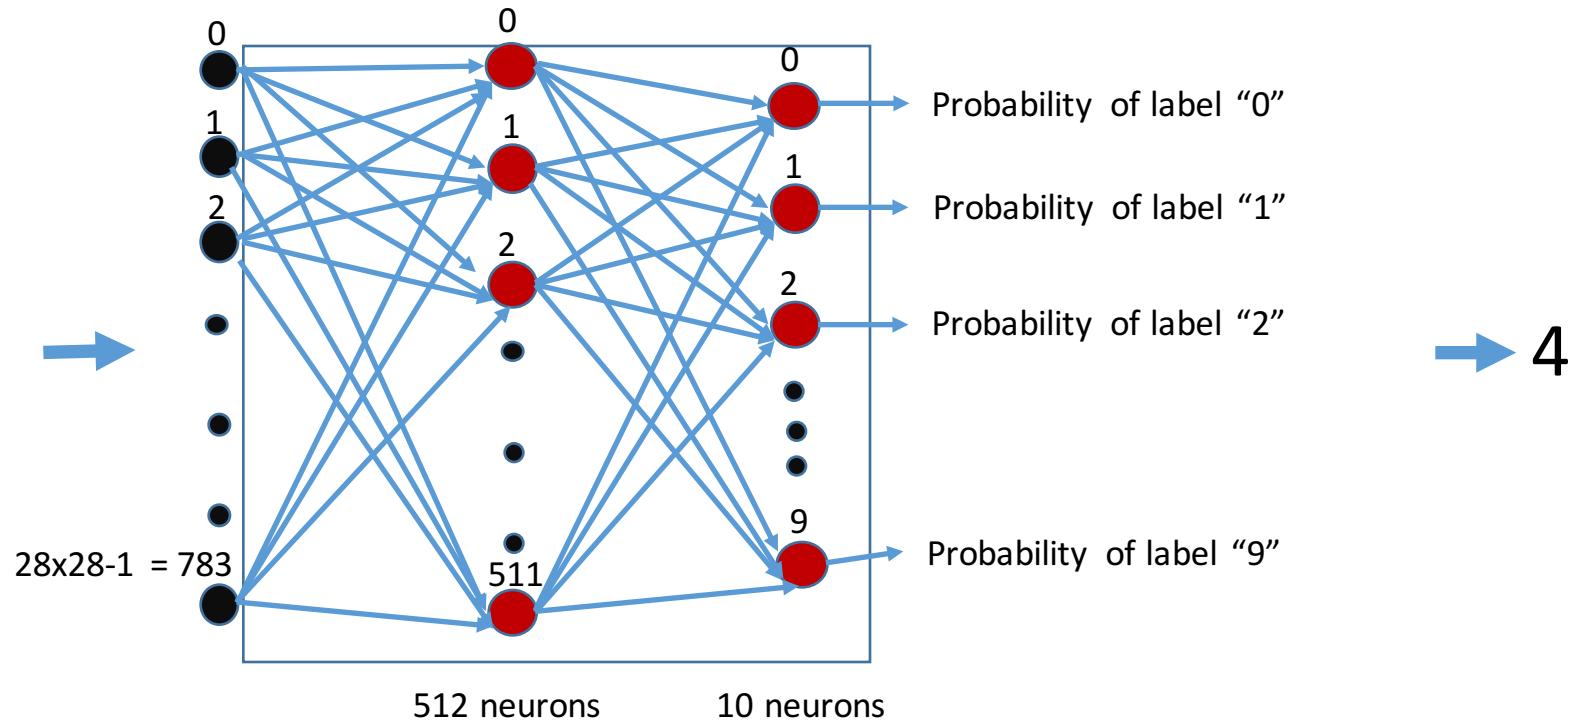
Loss function: categorical cross-entropy

Optimizer: RMSProp

Target Metric: Accuracy



28 x 28
2-d array



Step 4: Prepare training and test data

Here: Reshape and normalize input training data

Listing 2.4 Preparing the image data

```
train_images = train_images.reshape((60000, 28 * 28))  
train_images = train_images.astype('float32') / 255  
  
test_images = test_images.reshape((10000, 28 * 28))  
test_images = test_images.astype('float32') / 255
```

train_images:

Originally: 3-dimensional array of size $60000 \times 28 \times 28$, where each element is an integer in [0,255]

After reshaping: 2-dimensional array of size 60000×784 , where each element is an integer in [0,255]

After normalization: 2-dimensional array of size 60000×784 , where each element is a real number in [0,1]



28 x 28
2-d array

→ 1-d array of length
 $28 \times 28 = 784$

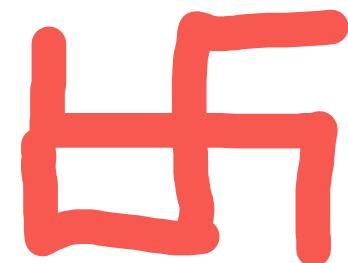
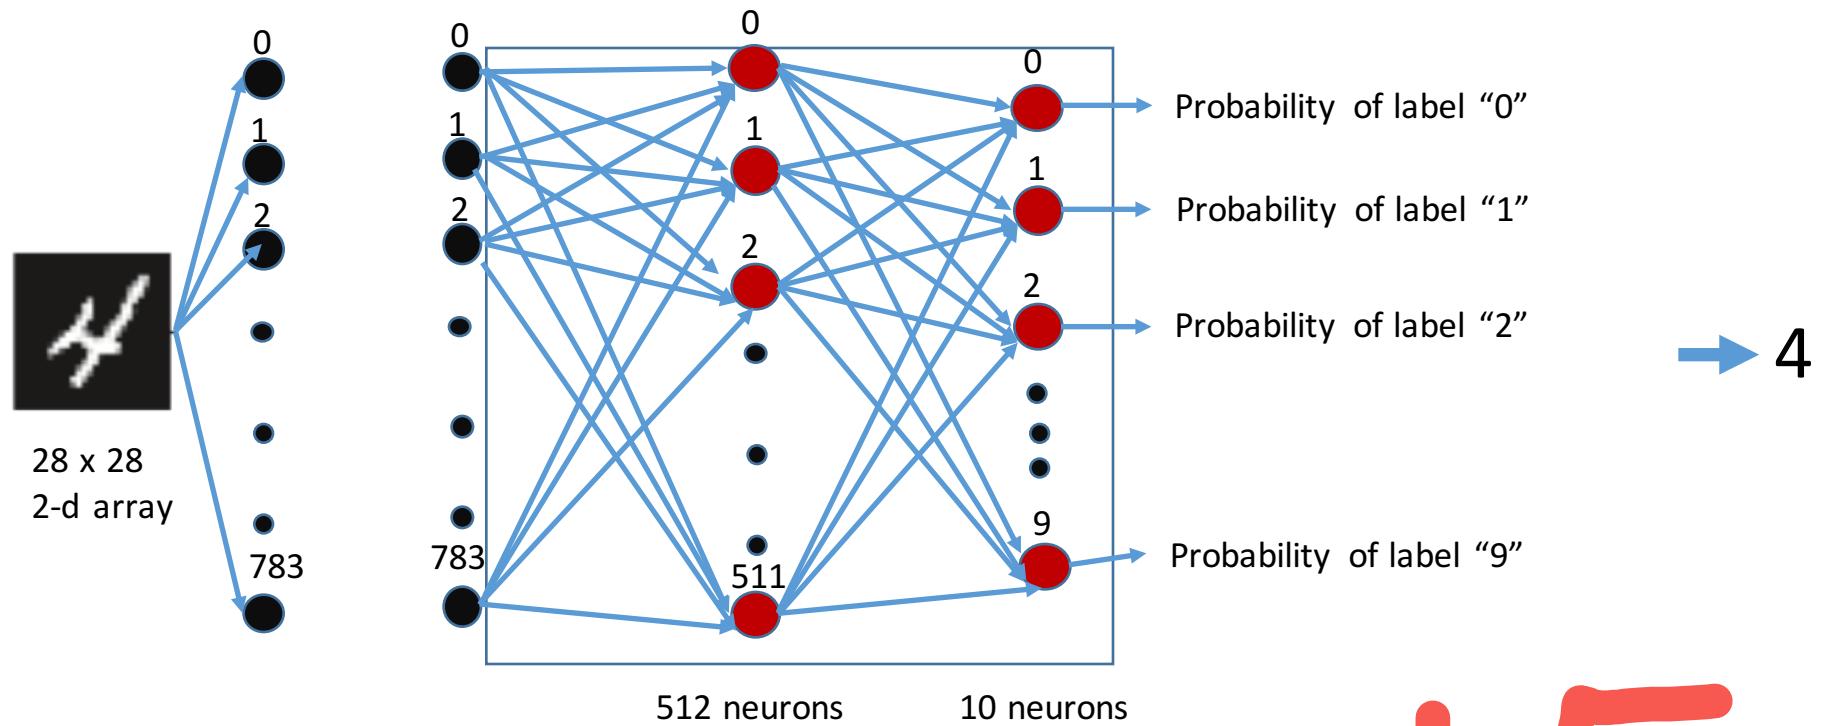
→ Normalize values in the array to between 0 and 1

The “**Teacher**”:

Loss function: categorical cross-entropy

Optimizer: RMSProp

Target Metric: Accuracy



“Reshape” output training data: categorically encode each label using one-hot encoding

```
from keras.utils import to_categorical  
  
train_labels = to_categorical(train_labels)  
test_labels = to_categorical(test_labels)
```

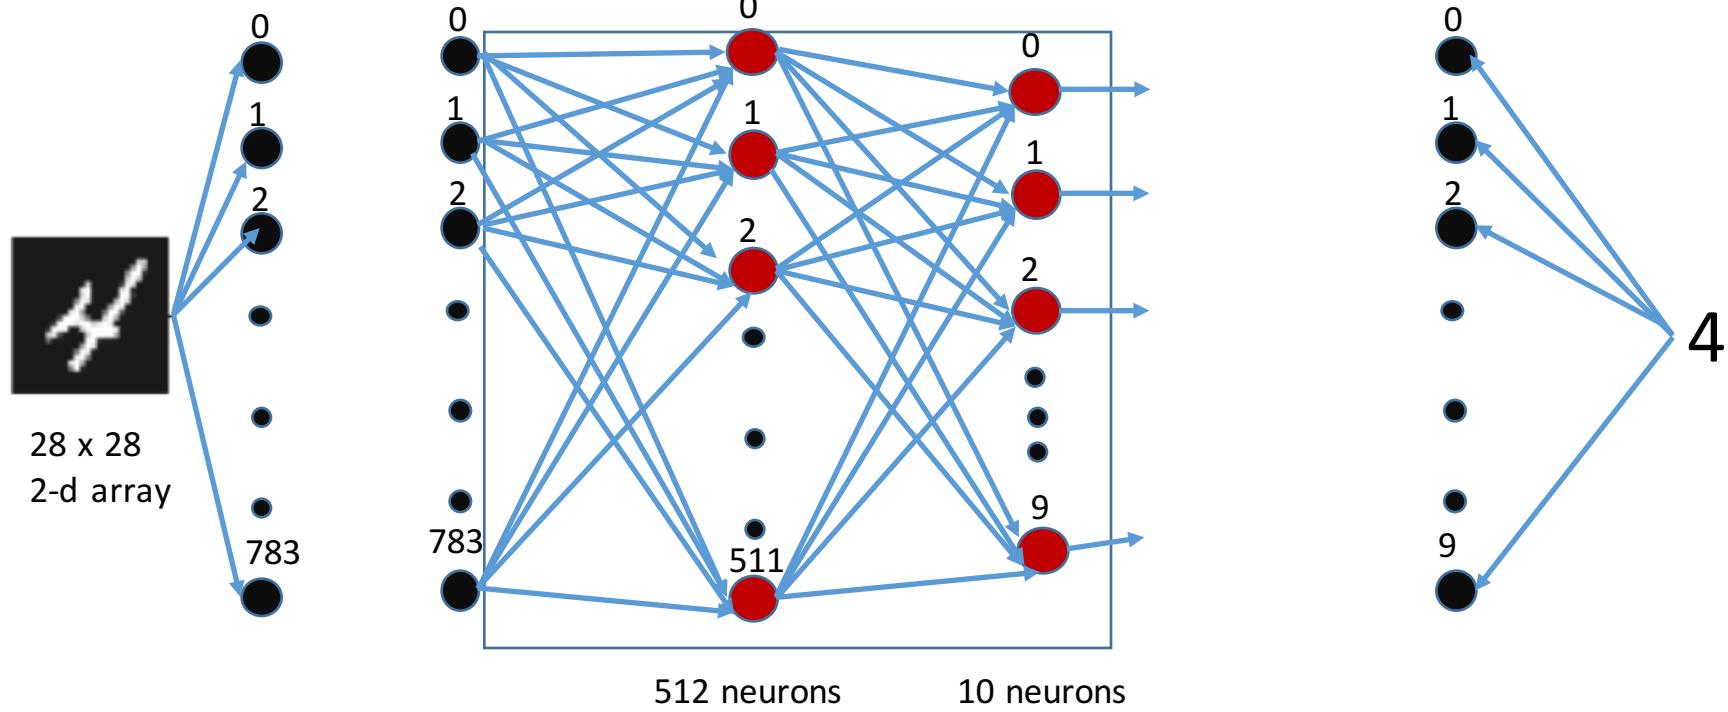
Label	One-hot encoding	Label	One-hot encoding
0	1,0,0,0,0,0,0,0,0,0	5	0,0,0,0,0,1,0,0,0,0
1	0,1,0,0,0,0,0,0,0,0	6	0,0,0,0,0,0,1,0,0,0
2	0,0,1,0,0,0,0,0,0,0	7	0,0,0,0,0,0,0,1,0,0
3	0,0,0,1,0,0,0,0,0,0	8	0,0,0,0,0,0,0,0,1,0
4	0,0,0,0,1,0,0,0,0,0	9	0,0,0,0,0,0,0,0,0,1

The “**Teacher**”:

Loss function: categorical cross-entropy

Optimizer: RMSProp

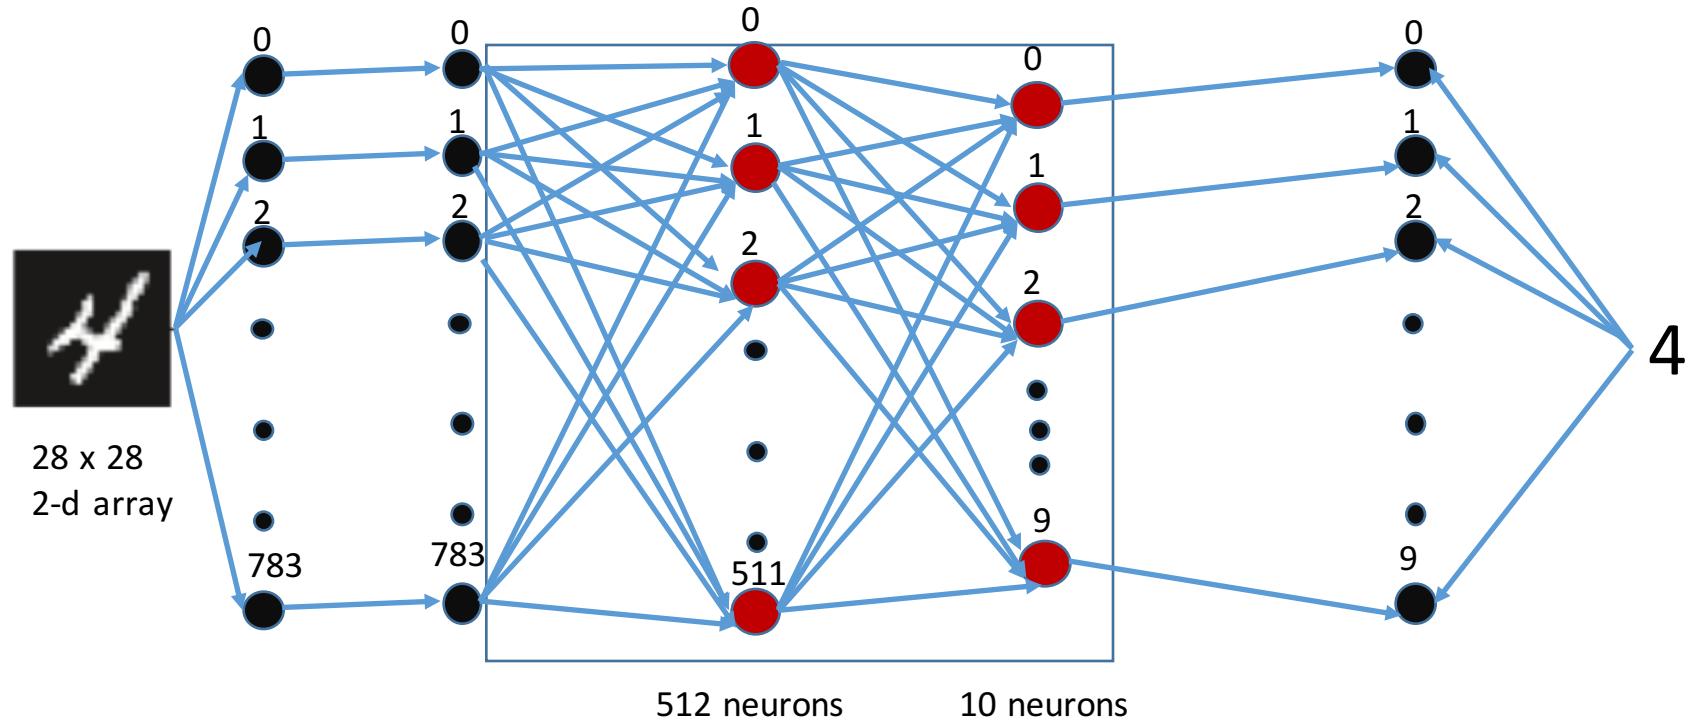
Target Metric: Accuracy



Step 5: Train the neural network

```
network.fit(train_images, train_labels, epochs=5, batch_size=128)
```

The “**Teacher**”: Loss function, Optimizer, Target Metric: Accuracy



Batch size: the number of samples to use each time for computing the loss function and updating the weights.

Epochs: the number of times the training process uses the whole training data set.

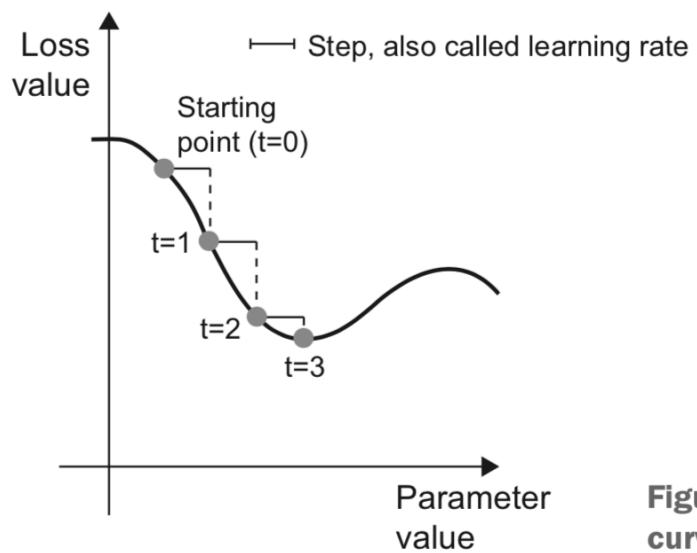


Figure 2.11 SGD down a 1D loss curve (one learnable parameter)

```
>>> network.fit(train_images, train_labels, epochs=5, batch_size=128)
Epoch 1/5
60000/60000 [=====] - 9s - loss: 0.2524 - acc: 0.9273
Epoch 2/5
51328/60000 [=====>.....] - ETA: 1s - loss: 0.1035 - acc: 0.9692
```

And so on (totally 5 epochs).

Accuracy on training data: 97.8%

Step 6: Test the trained neural network

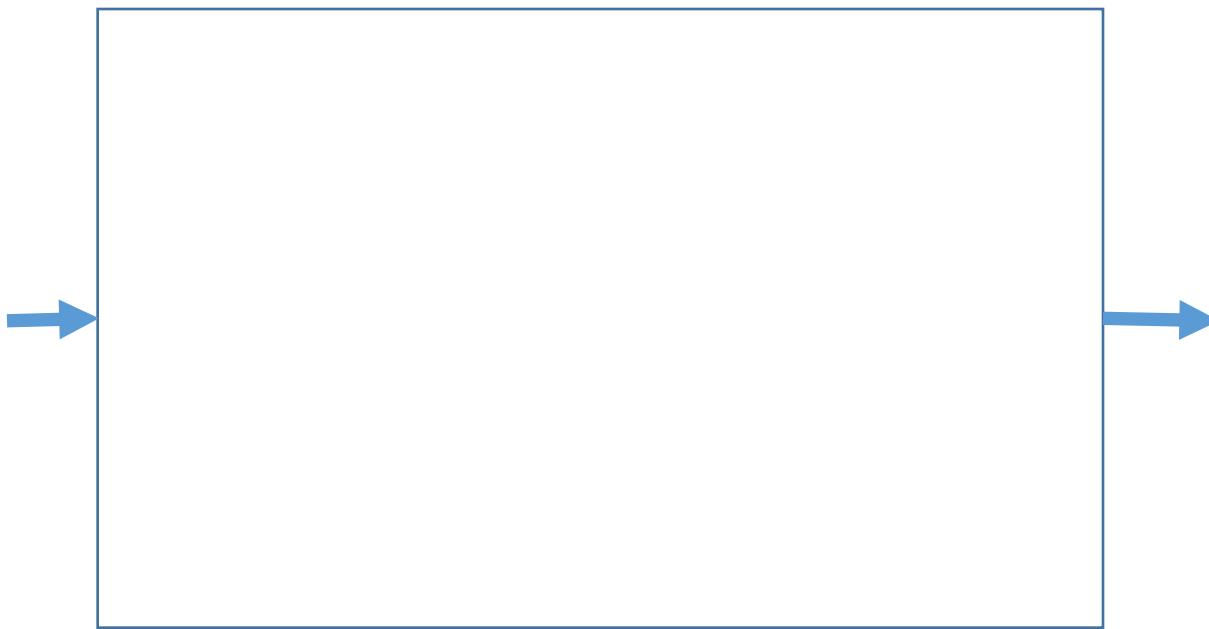
```
>>> test_loss, test_acc = network.evaluate(test_images, test_labels)
>>> print('test_acc:', test_acc)
test_acc: 0.9785
```

Compare to training accuracy: 0.989

Test accuracy is (clearly) lower than training accuracy.
Maybe there is some over-fitting to data.

But still, performance is nice!

Summary



Step 1: Load the dataset

Listing 2.1 Loading the MNIST dataset in Keras

```
from keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

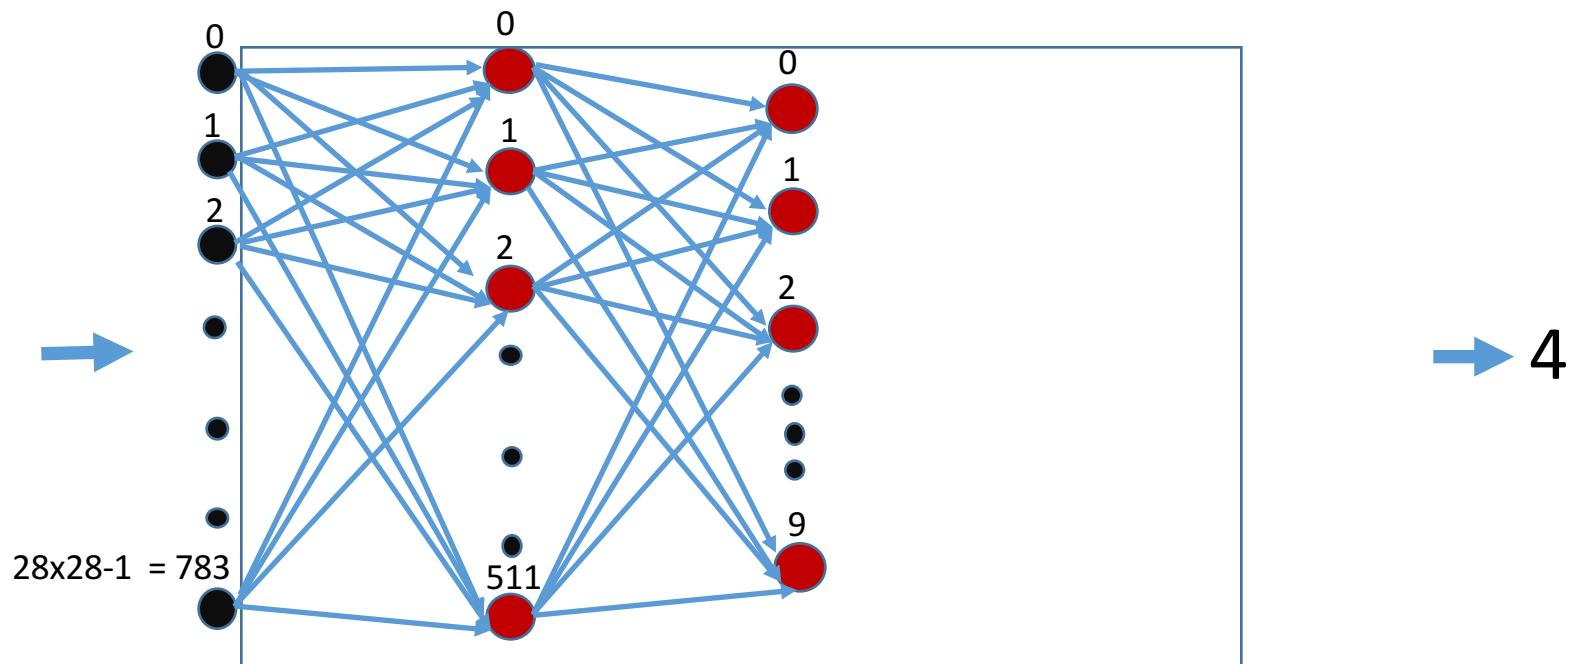
A large, empty rectangular box with a blue border, centered on the page.
4

Step 2: Build neural network architecture

```
from keras import models  
from keras import layers  
  
network = models.Sequential()  
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28, )))  
network.add(layers.Dense(10, activation='softmax'))
```



28 x 28
2-d array



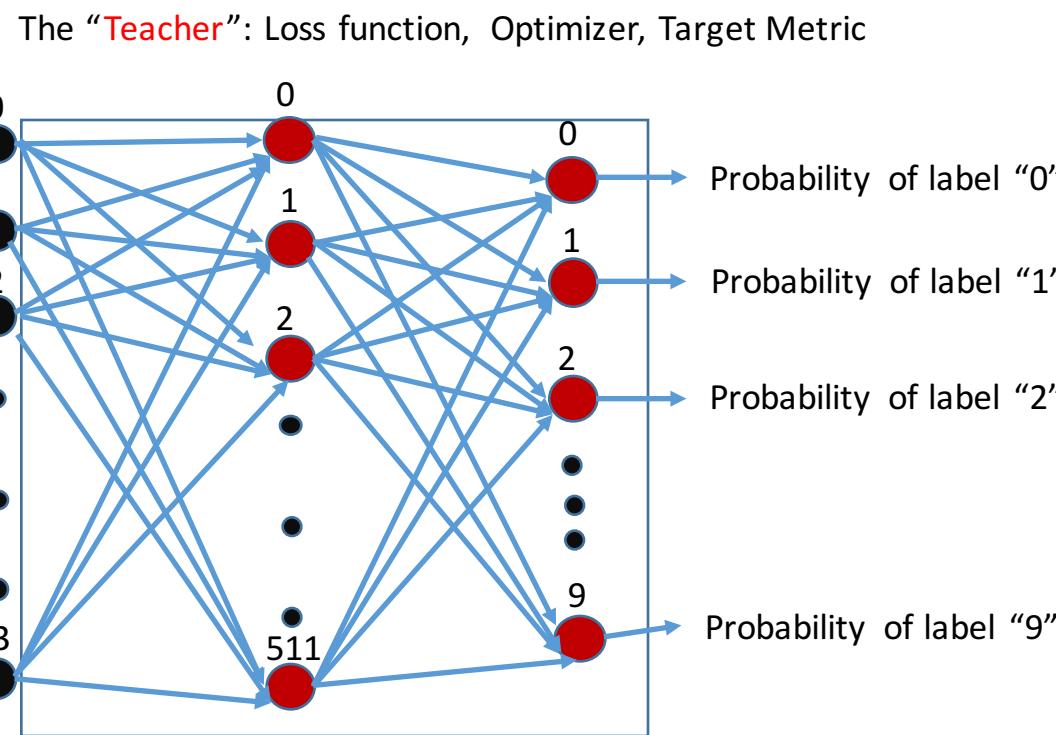
Step 3: choose loss function, optimizer, and target metrics

```
network.compile(optimizer='rmsprop',  
                 loss='categorical_crossentropy',  
                 metrics=['accuracy'])
```



28 x 28
2-d array

$$28 \times 28 - 1 = 783$$

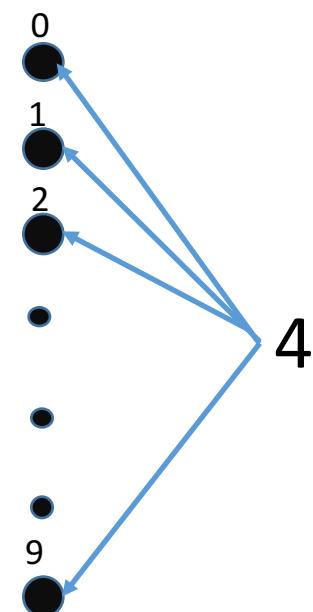
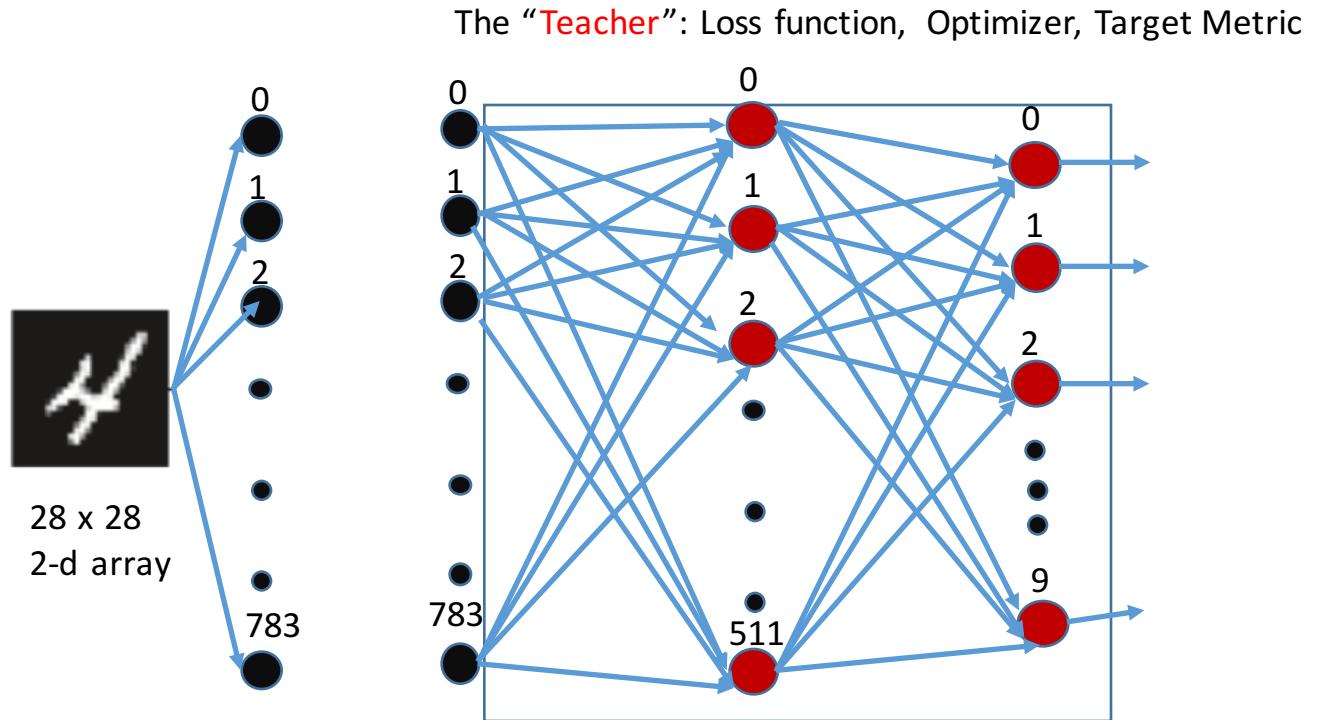


→ 4

Step 4: Prepare training and test data

```
train_images = train_images.reshape((60000, 28 * 28))  
train_images = train_images.astype('float32') / 255  
  
test_images = test_images.reshape((10000, 28 * 28))  
test_images = test_images.astype('float32') / 255
```

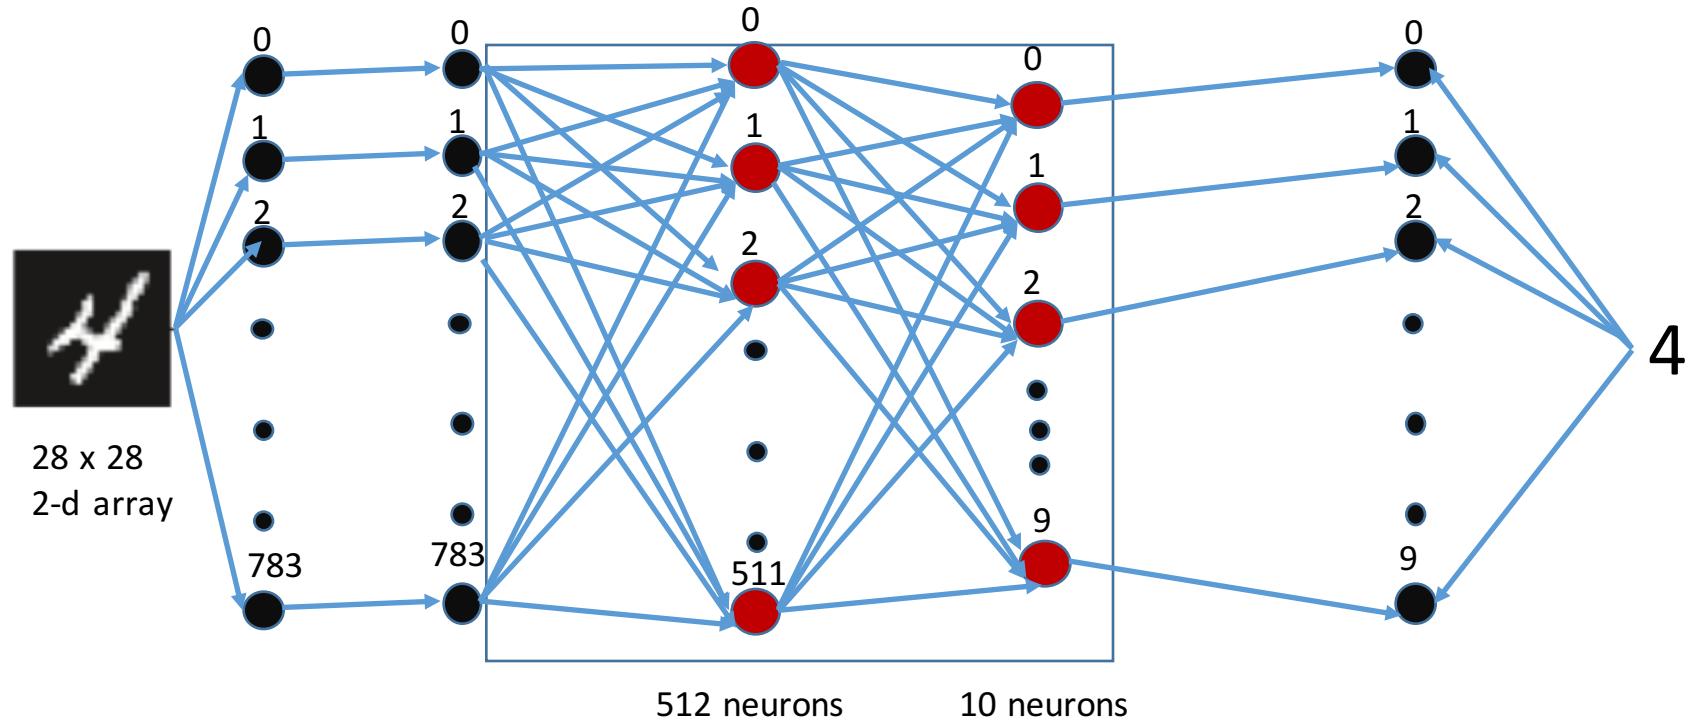
```
from keras.utils import to_categorical  
  
train_labels = to_categorical(train_labels)  
test_labels = to_categorical(test_labels)
```



Step 5: Train the neural network

```
network.fit(train_images, train_labels, epochs=5, batch_size=128)
```

The “**Teacher**”: Loss function, Optimizer, Target Metric: Accuracy



Step 6: Test the trained neural network

```
>>> test_loss, test_acc = network.evaluate(test_images, test_labels)
>>> print('test_acc:', test_acc)
test_acc: 0.9785
```

How did I do?

Well ...



Miscellaneous Basic Concepts

Data representation: Tensor (Array)

- Scalar numbers (0-dimensional tensors)
- Vectors (1-d tensors)
- Matrices (2-d tensors)
- 3-d tensors, and higher-dimensional tensors
- Key attributes for a tensor:
 - (1) number of axes
 - (2) **shape**
 - (3) data type

Some basic tensor operations

- Add two tensors (of the same shape): element-wise addition
- Apply a ReLU activation function to a tensor: element-wise operation
- **Tensor Product** (also called tensor dot)

$$A_5^T B_5 = C \text{ (scalar number)}$$

$$A_{2,3,4,5} B_5 = C_{2,3,4}$$

$$A_{7,5} B_5 = C_7$$

$$A_{2,3,4,5} B_{5,6} = C_{2,3,4,6}$$

$$A_7^T B_{7,5} = C_5^T$$

$$A_{2,3,4,5} B_{5,6,7} = C_{2,3,4,6,7}$$

$$A_{2,8} B_{8,6} = C_{2,6}$$

$$A_{2,3,4,5} B_{5,4,7} = C_{2,3,4,4,7}$$

Reshape tensor

```
>>> x = np.array([[0., 1.],  
                 [2., 3.],  
                 [4., 5.]])  
>>> print(x.shape)  
(3, 2)
```

```
>>> x = x.reshape((6, 1))  
>>> x  
array([[ 0.],  
       [ 1.],  
       [ 2.],  
       [ 3.],  
       [ 4.],  
       [ 5.]])  
  
>>> x = x.reshape((2, 3))  
>>> x  
array([[ 0.,  1.,  2.],  
       [ 3.,  4.,  5.]])
```

Basic terms for a neural network

- Layers: the building blocks in a neural network
- Model: network of layers
- Loss function and optimizer: keys to configuring the learning process

Keras: a deep learning library for Python

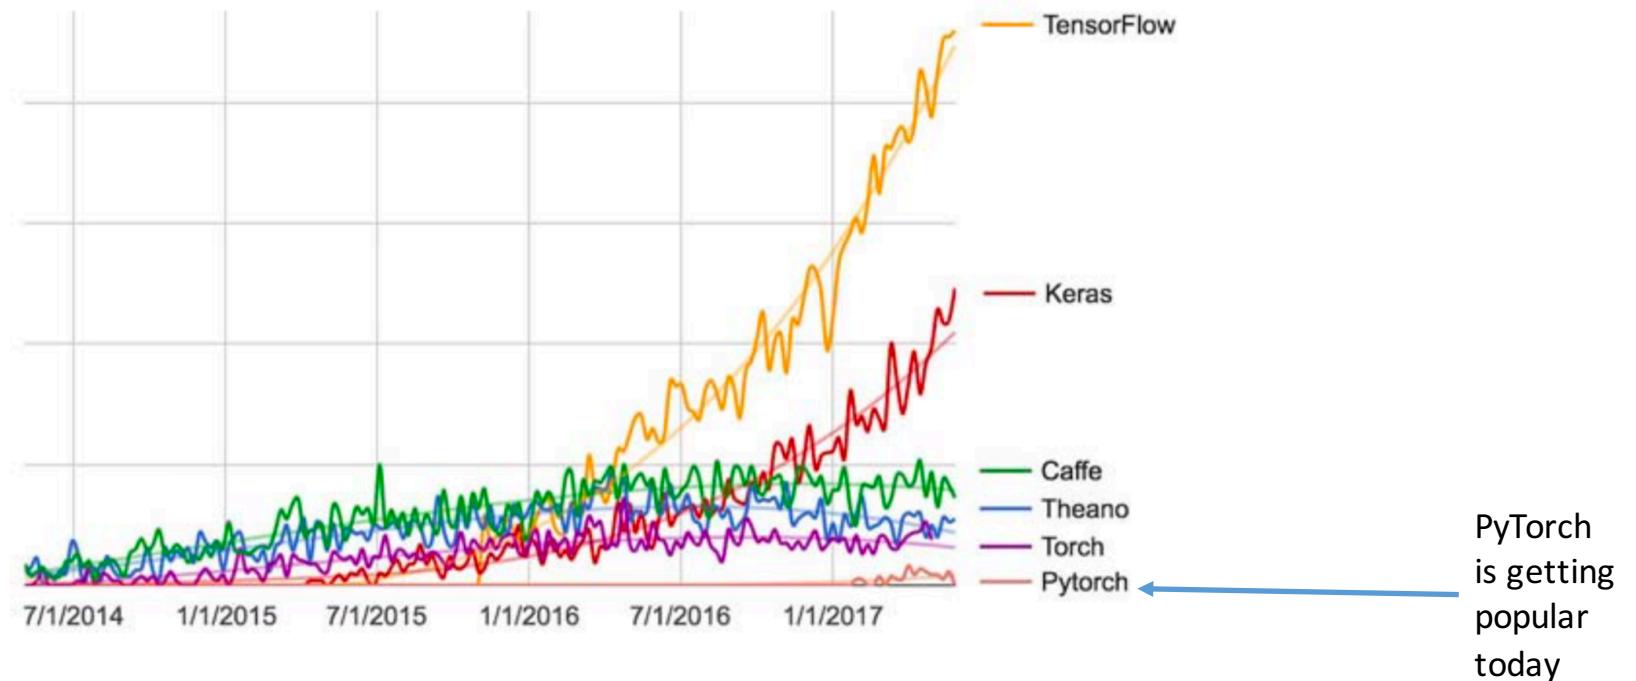


Figure 3.2 Google web search interest for different deep-learning frameworks over time

Keras: a deep learning library for Python

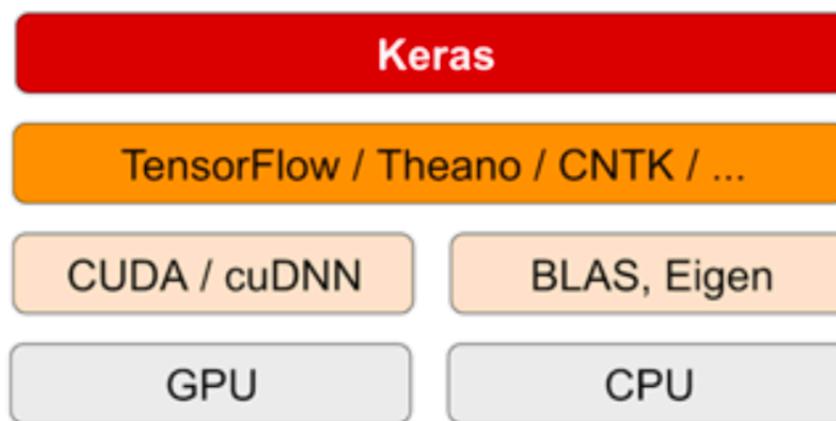


Figure 3.3 The deep-learning software and hardware stack

Use a GPU when possible

Jupyter notebook: a nice way to edit and run deep learning experiments

